

# Kubernetes Technical Primer



Ben Coleman

@BenCodeGeek

*Feb 2020, v1.4*

## Content Notes



This guide is not intended to be delivered as a complete presentation

It can be used as a learning guide or hand out

Alternatively take out sections or slides and use/present as needed

## Target Audience



Technical architects, developers and platform engineers

Those wanting to learn the core fundamentals of Kubernetes

Not for those trying to deploy Kubernetes from scratch or get into internals

## [Introduction To Kubernetes](#)

What is Cloud Native?  
The need for orchestration  
Kubernetes: the industry leading orchestrator  
The elements of orchestration  
Why Choose Kubernetes?  
Kubernetes – A Modern Orchestrator  
Core Concepts and Terms  
Simplified Architecture  
Internal Architecture  
Highly Extensible

## [Core Components](#)

Pods  
Deployments & Replica Sets  
Services  
Services – Illustrated  
DNS and Service Discovery  
Data Volumes & Mounts  
Persistent Volumes  
Stateful Sets  
Daemon Sets  
Jobs & CronJobs  
Namespaces  
Role Based Access Control (RBAC)  
Putting It All Together

## [Using Kubernetes](#)

Command Line - kubectl  
Kubectl – Common Commands  
Dashboard  
Accessing the Dashboard  
Kubernetes Object Management  
Introduction to the Declarative Model  
Idempotent Updates & Desired State  
Labels & Selectors

### [Configuring Basic Workloads](#)

Environmental Variables  
Secrets  
ConfigMaps  
Resource Management  
Liveness & Readiness Probes  
Commands & Arguments

## [Beyond The Basics](#)

Service Mesh  
Kubernetes Ecosystem – Common Projects  
The Kubernetes API

### [Additional Network Services](#)

Ingress  
External DNS  
Cert Manager  
Putting It All Together

### [Debugging and Troubleshooting Workloads](#)

Describing Objects  
Container Logs  
Get Shell Access

### [Advanced Pod Configuration](#)

Deeper Dive on Manifests  
Init Containers  
Node Selector  
Affinity and Taints  
Sidecars

### [Scaling](#)

Manually Scaling  
Horizontal Pod Autoscaler (HPA)  
Cluster Autoscaler (CA)

## [DevOps](#)

Modern Infrastructure  
DevOps Containers Lifecycle Loop  
Cluster Isolation Patterns: Physical Isolation  
Cluster Isolation Patterns: Logical Isolation  
Helm Introduction  
Helm – The Basics

## [Azure Kubernetes Service](#)

Managed Kubernetes on Azure – AKS  
AKS Virtual Nodes  
AKS Networking Models  
AKS Advanced Networking  
Azure Active Directory  
'HTTP application routing' Add-On  
'Monitoring' Add-On  
Azure DevSpaces

## [Additional Resources](#)

# MAIN SECTIONS

Introduction To  
Kubernetes



Core Components



Using Kubernetes



Beyond The Basics



DevOps



Azure Kubernetes  
Service



# Introduction To Kubernetes



# What is Cloud Native?



“Cloud native technologies empower organizations to build and run scalable applications in **modern, dynamic environments** such as public, private, and hybrid clouds. **Containers, service meshes, microservices**, immutable infrastructure, and declarative APIs exemplify this approach.

These techniques enable **loosely coupled systems** that are **resilient, manageable, and observable**. Combined with **robust automation**, they allow engineers to make high-impact changes frequently and predictably with minimal toil.”



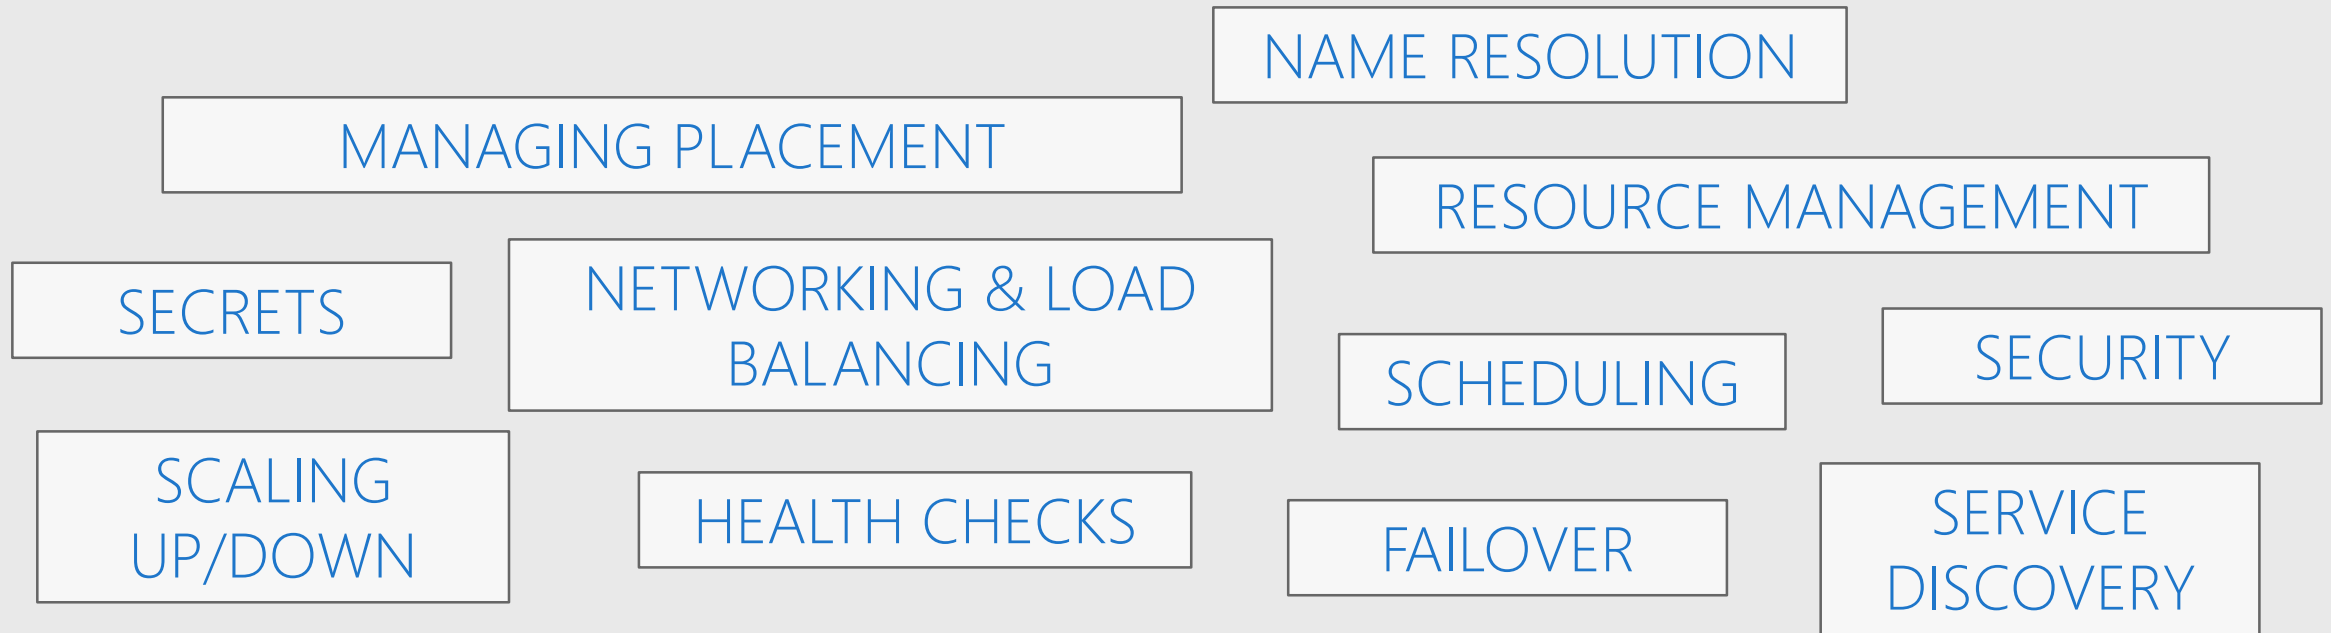
**CLOUD NATIVE**  
**COMPUTING FOUNDATION**

# The need for orchestration



OK great,  
But now what?

# The need for orchestration



Production And At Scale



# Kubernetes - production grade orchestration



## Portable

Public, private, hybrid,  
multi-cloud

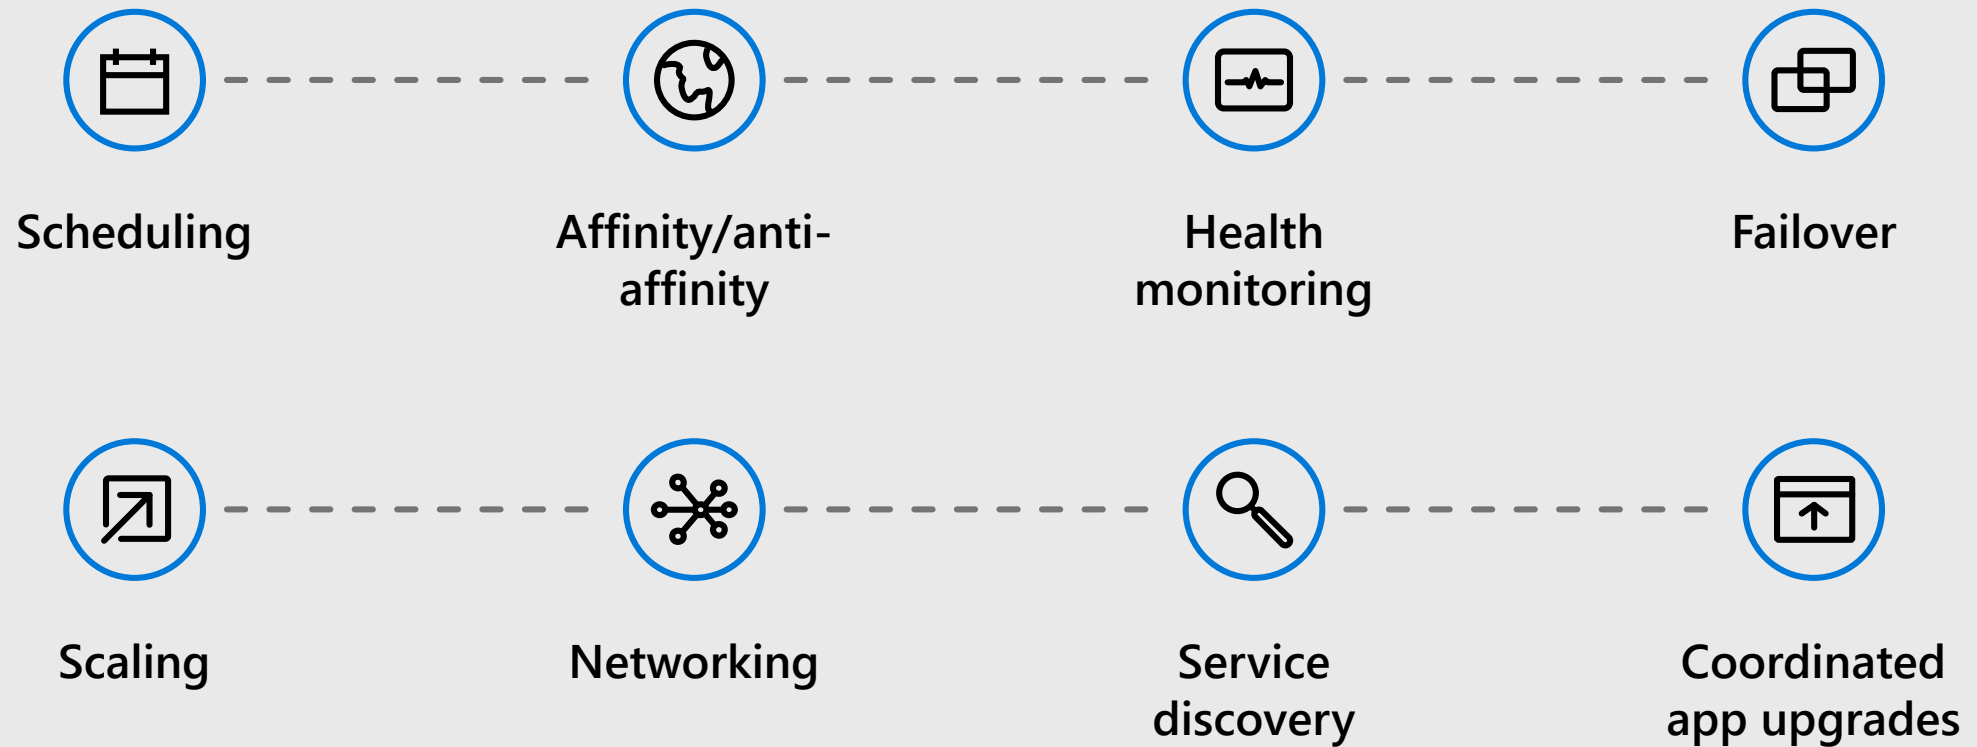
## Extensible

Modular, pluggable,  
hookable, composable

## Self-healing

Auto-placement, auto-restart,  
auto-replication, auto-scaling

# The elements of orchestration



# Why Choose Kubernetes?



Cornerstone of cloud native approach



Run anywhere



Industry adoption



Open Source with high degree of support



Avoids lock-in



Skills availability

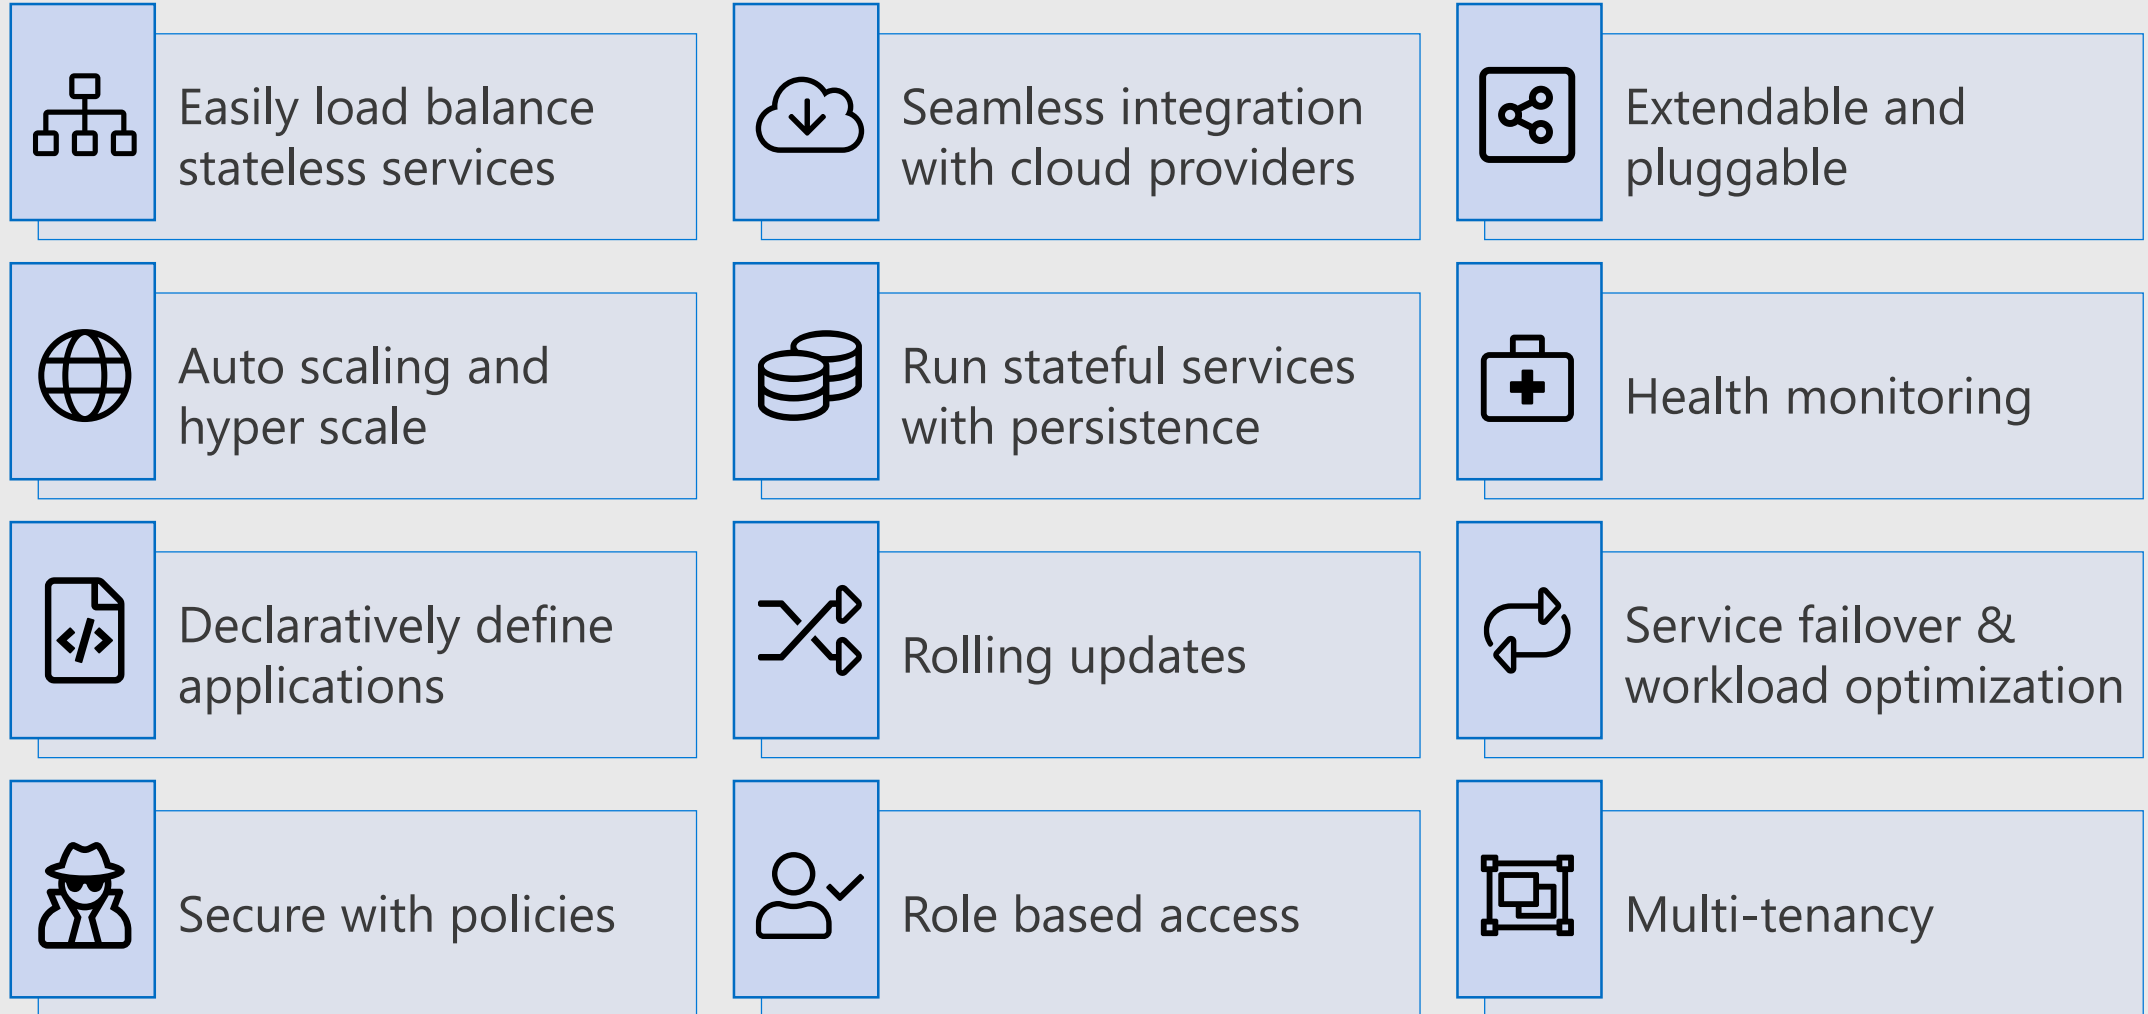


Large & growing ecosystem




Rise of microservices & containers

# Kubernetes – A Modern Orchestrator



# Core Concepts and Terms



## Node

A worker machine (VM) normally clustered, each capable of running pods



## Deployment

A logical object for managing a replicated application (i.e. set of pods)



## Label

Metadata attached to any object for configuration and selection



## Pod

A group of one or more containers that is lifecycle managed



## Service

Network access to a resource, e.g. pod or port. Typically load balanced

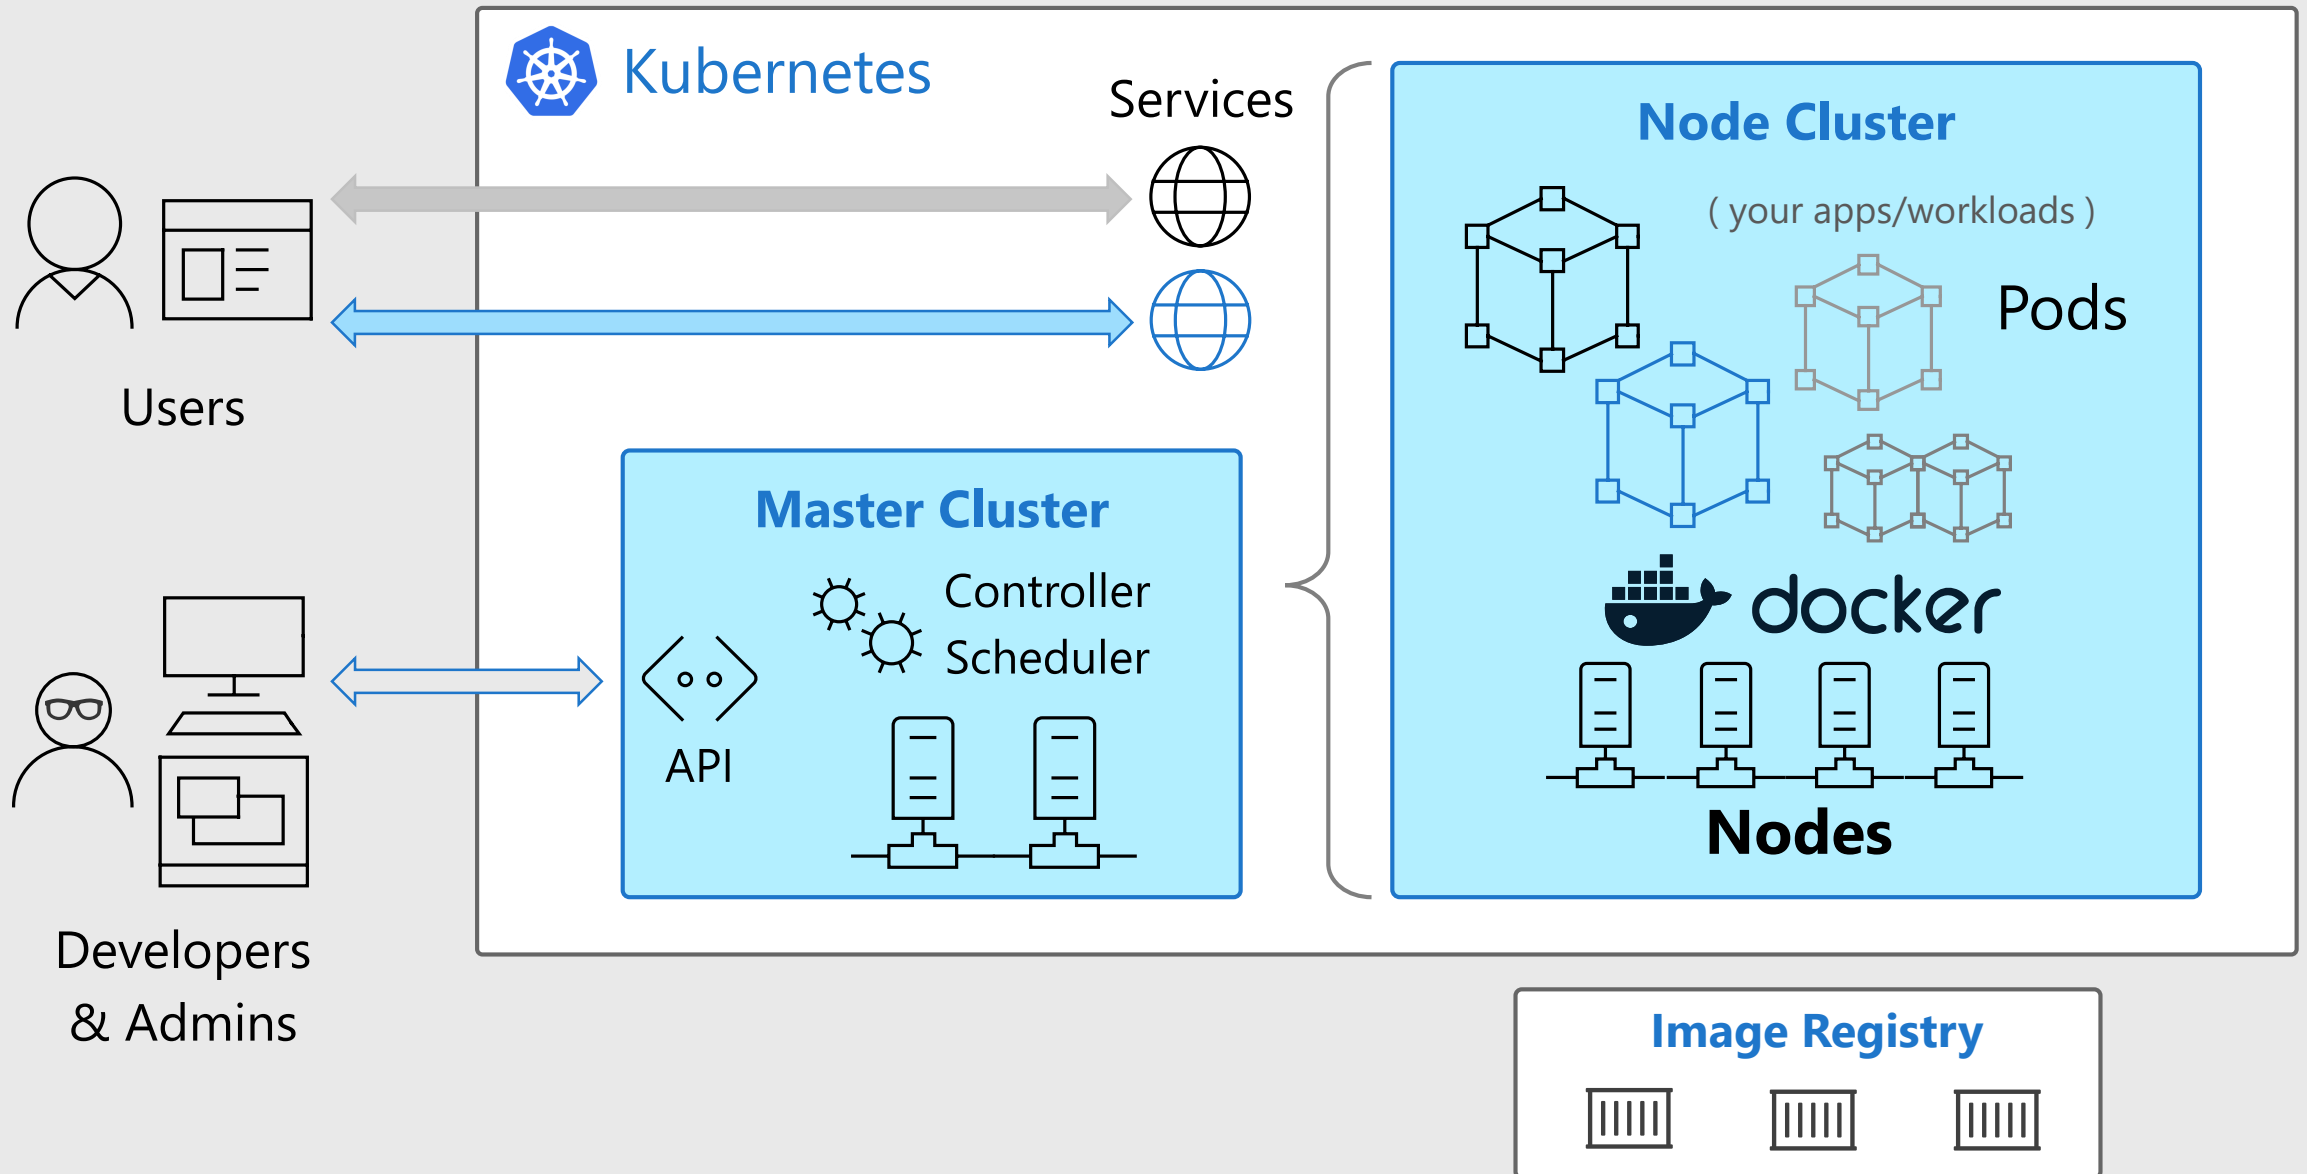


## Replica Set

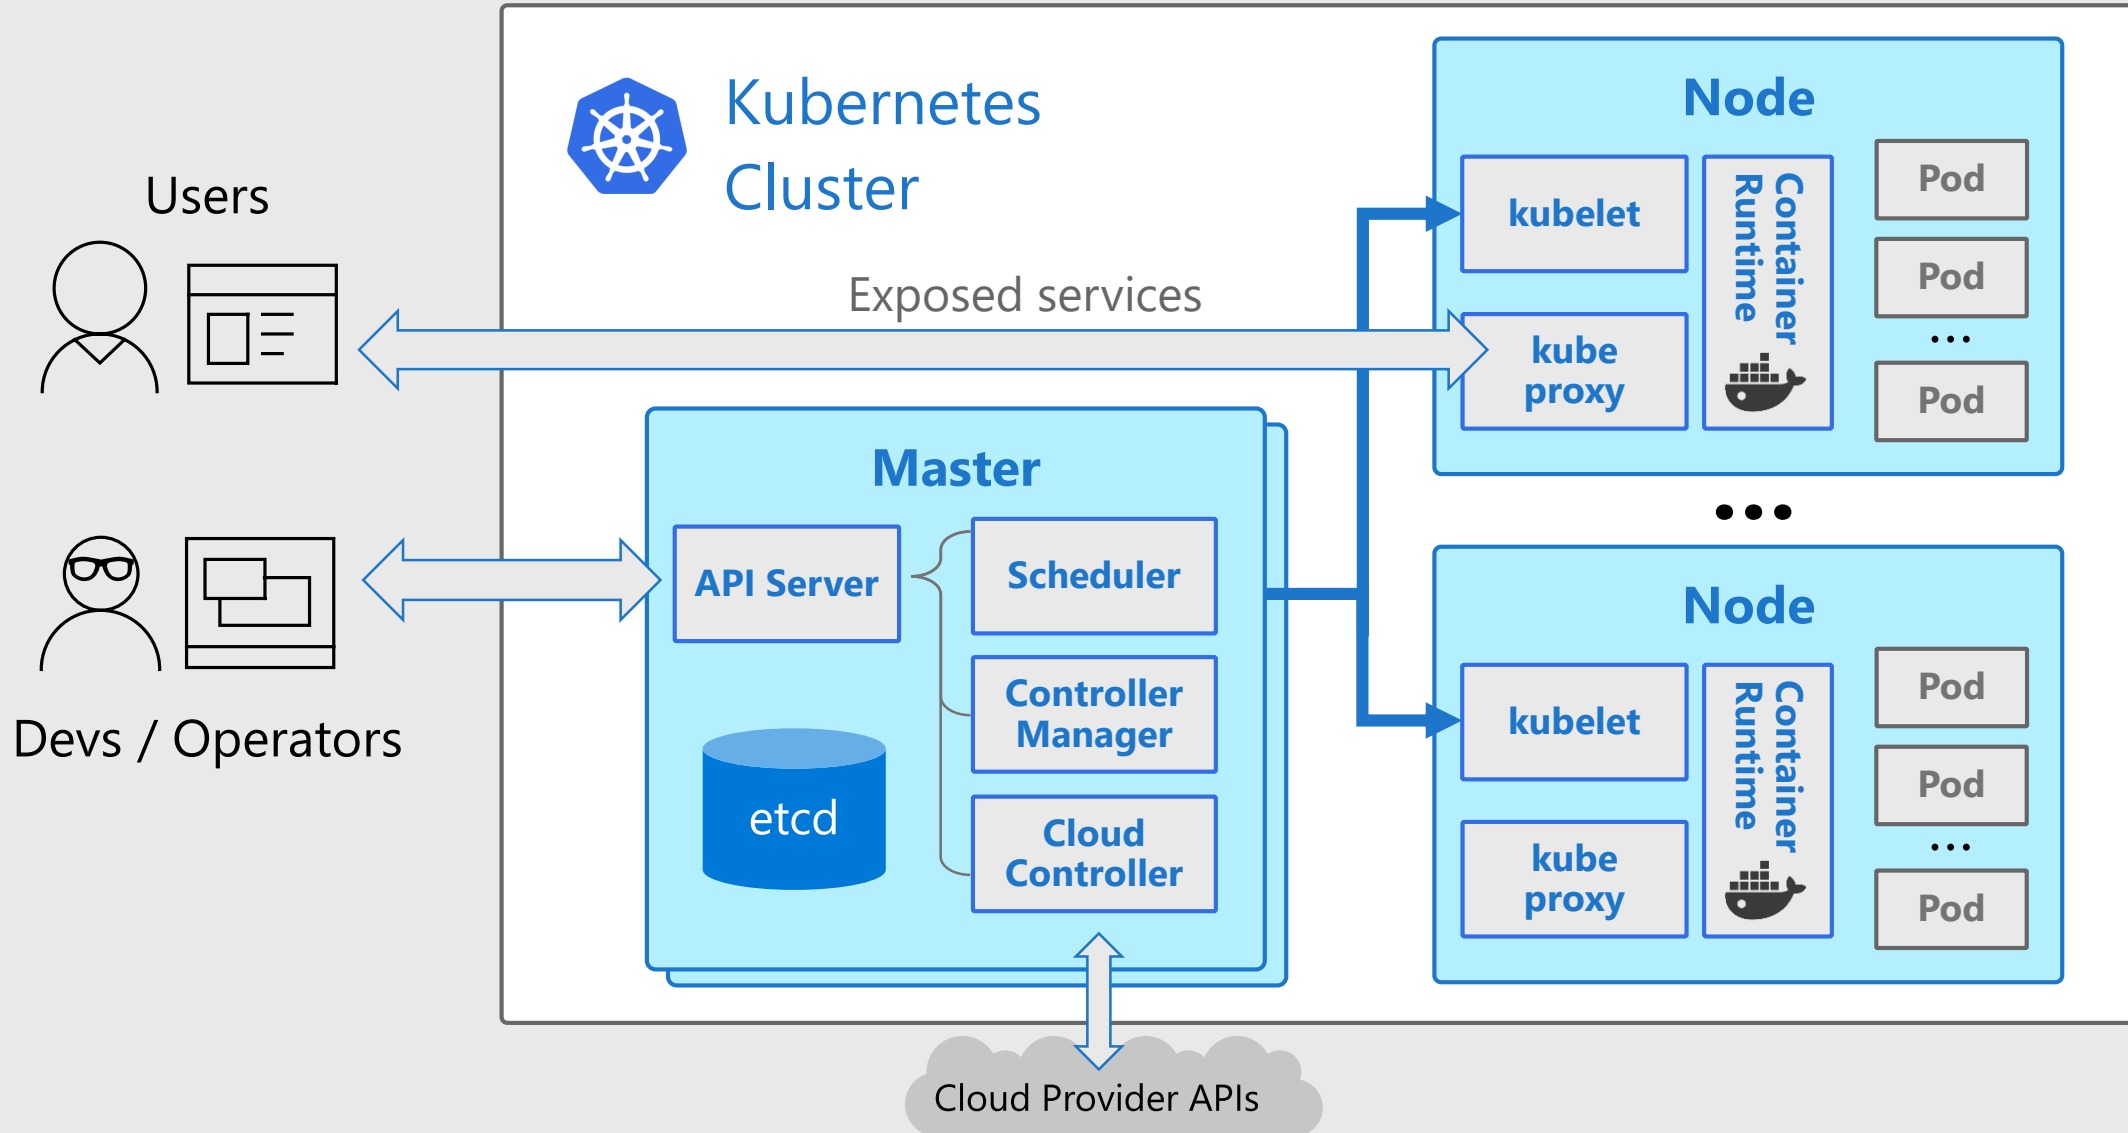
A set of one or more pods that is distributed and replicated across Nodes



# Simplified Architecture

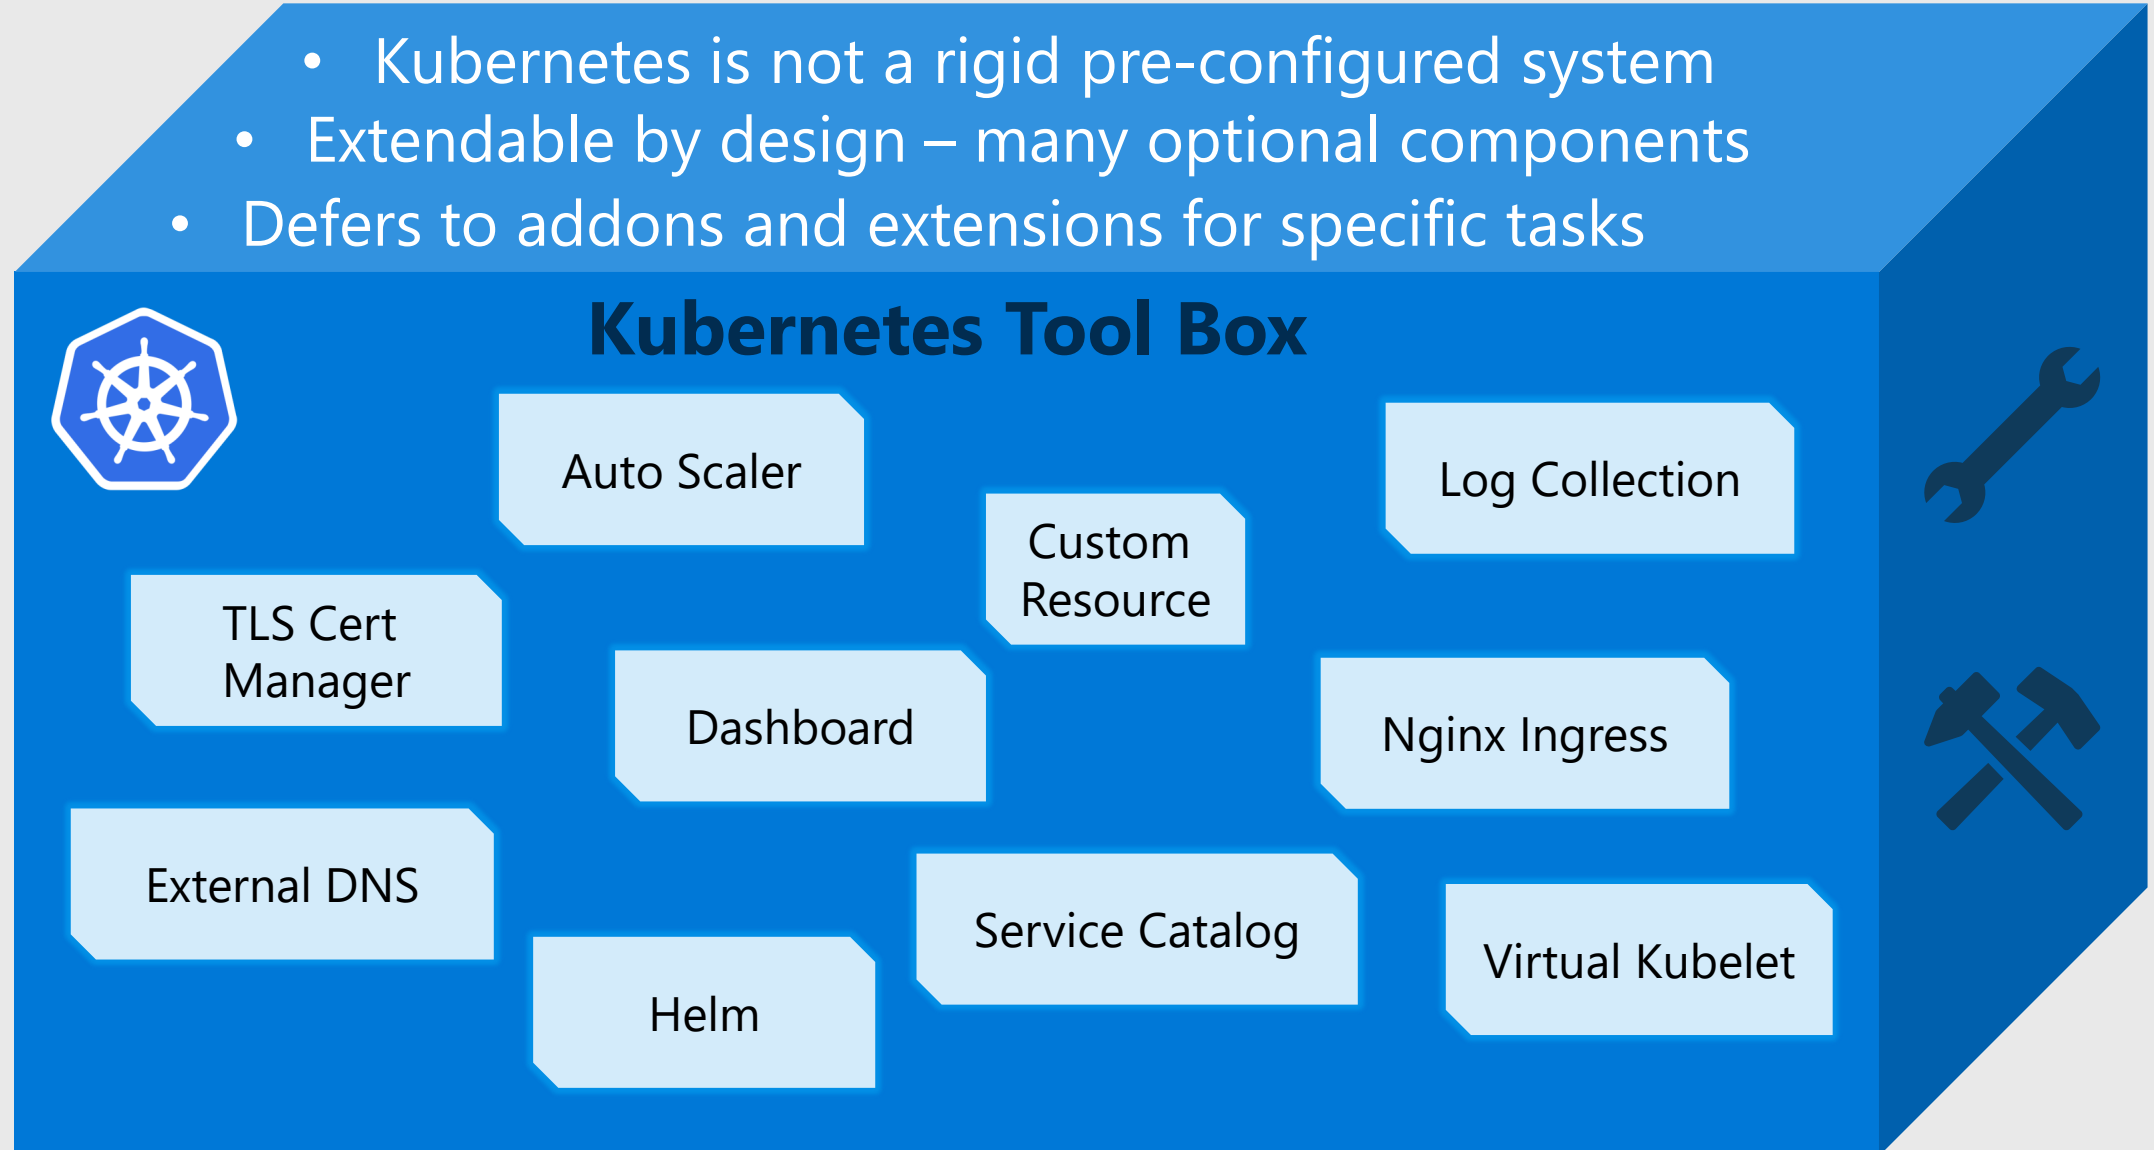


# Internal Architecture



# Highly Extensible / Unopinionated

- Kubernetes is not a rigid pre-configured system
- Extendable by design – many optional components
- Defers to addons and extensions for specific tasks





# Core Components



# Pods

## Fundamental building block of Kubernetes


*Pods* run **one or more** containers

Containers in a pod **share network/storage**

*Pods* each have their own IP address

*Pods* expose **one or more** ports

*Pods* are scheduled and run on a *Node*



*Example Pod*

**Pods are the primary way of running your workloads in Kubernetes**

# Deployments & Replica Sets

Scale and run pods across multiple nodes

*Deployments* describe a **replicated** set of *Pods*

A *Deployment* represents **desired state**

- Rolling updates used to safely roll out changes

You **scale** *Deployments* up & down


*Deployments* typically run **stateless** workloads

*Deployments* use *ReplicaSets*

```
Deployment
- name: MyApp
- replicas: 3
```



```
ReplicaSet
- replicas: 3
```



Deployments let you run & scale stateless workloads in Kubernetes

# Services


## Network access to Pods

*Pods* are ephemeral - they can move/die/change without notice. In general, users do not directly access *Pods*

*Services* are an abstraction which defines a **logical set** of *Pods* and a policy by which to access them

*Services* use labels and selectors to map to *Pods*

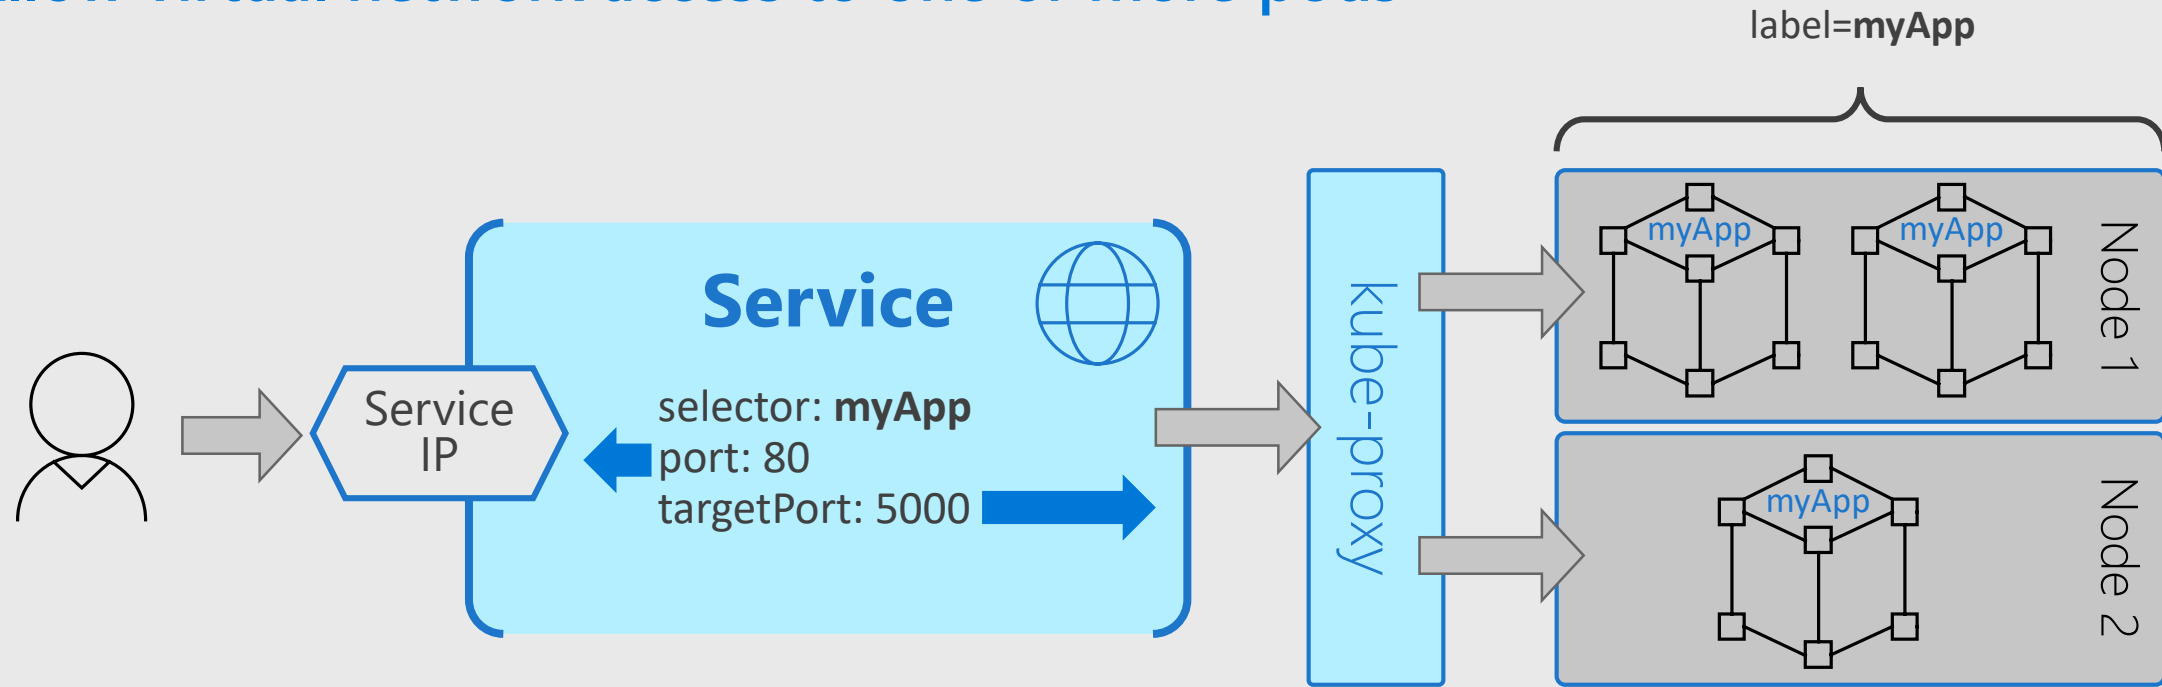
*Services* are assigned IP addresses and DNS names



Services are how you connect to Pods over the network

# Services – Illustrated

Allow virtual network access to one or more pods



EXTERNAL

## LoadBalancer

Uses cloud provider to present an external load-balanced IP

INTERNAL

## ClusterIP

Internal virtual IP, only accessible by other pods/services

*Note. Uses 'round robin' to select pods*

# DNS and Service Discovery

## Naming for Pods and Services


DNS in cluster is handled automatically by Kubernetes

All ClusterIP (internal) *Services* get assigned a DNS record based on the service's name

*Pods* also get DNS names, but this is less useful


External public DNS can also be configured – See “Additional Network Services” section later

```
$ nslookup myapp  
Address: 10.200.4.30
```



```
Service  
- name: myapp
```

Service IP: 10.200.4.30



Connect to your workloads using DNS and Services

# Data Volumes & Mounts


## Handling State & Data

Filesystem of running containers is ephemeral. All data written will be lost on a restart

Use **Volumes** to hold data or state you want to keep, or to inject data into a *Pod*

Volumes are mounted into a container at a **mountPath**

Many types of storage can be used to back the Volume



*Warning! A volume shares lifecycle with the Pod, so are not persistent*

Volumes hold data and state for Pods and containers

# Persistent Volumes

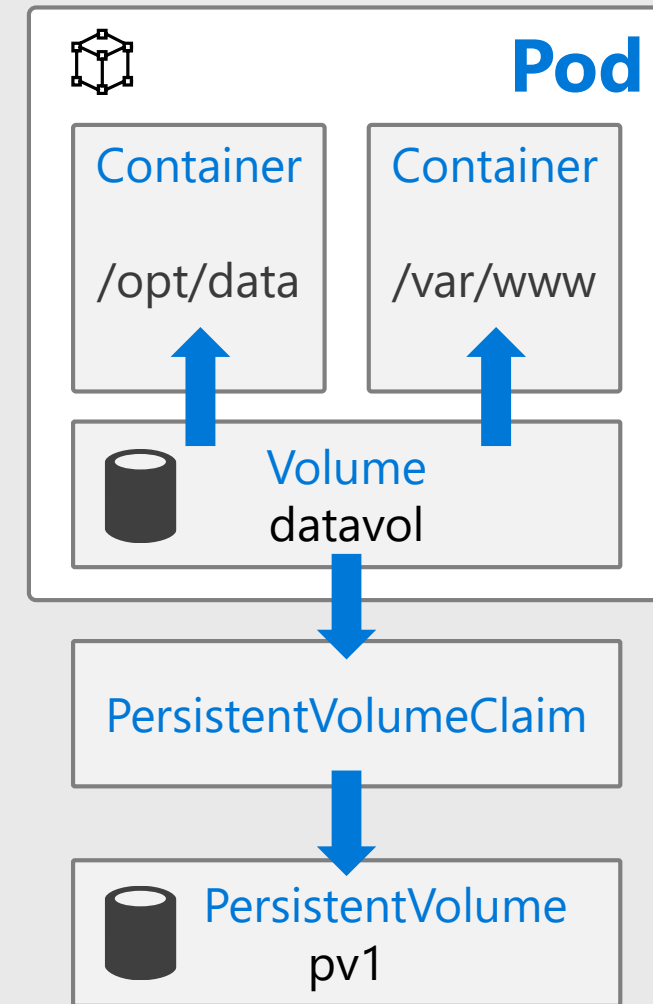
## Handling State & Data

A *PersistentVolume* allows you to hold data independent of Pod lifecycle

A pod uses a *PersistentVolumeClaim* to bind to a *PersistentVolume*

- **ReadWriteOnce** - mounted on a single Node (e.g. db)
- **ReadWriteMany** - mounted on multiple Nodes

Many **storage plugins** exist: NFS, iSCSI, Azure (Disk & Files), Flocker, Ceph, Gluster, AWS



**Persistent Volumes retain data long term, outside of Pods**



# Stateful Sets


## Handling Stateful workloads

A *StatefulSet* is like a *Deployment* except *Pods* get well defined names and replicas start in **ordered sequence**

*StatefulSets* **retain identity** regardless of which *Node* they run on

Each *Pod* in a *StatefulSet* will bind to the same defined *PersistentVolumeClaim*

```
StatefulSet
- name: MyDbSet
- serviceName: "mongo"
- replicas: 2
```



Use a StatefulSet rather than Deployment for stateful workloads

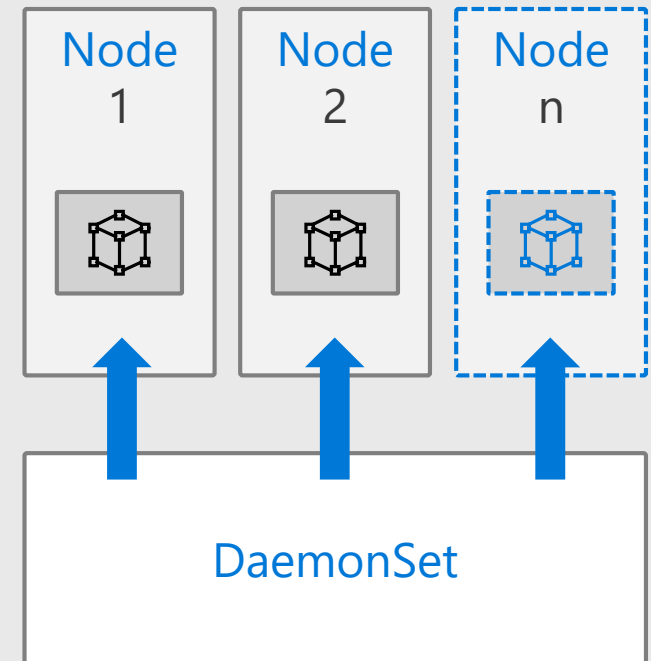
# Daemon Sets

## Running Pods across all Nodes

A *DaemonSet* ensures that all *Nodes* in the cluster run a given *Pod*. *Pods* will be **created/removed as *Nodes* are added/removed**

Used for special system and cluster daemons, logging, storage, etc.

*DaemonSets* are **not** often used for normal app workloads



DaemonSets run system Pods for monitoring & network

# Jobs & CronJobs

## Workloads that run to completion

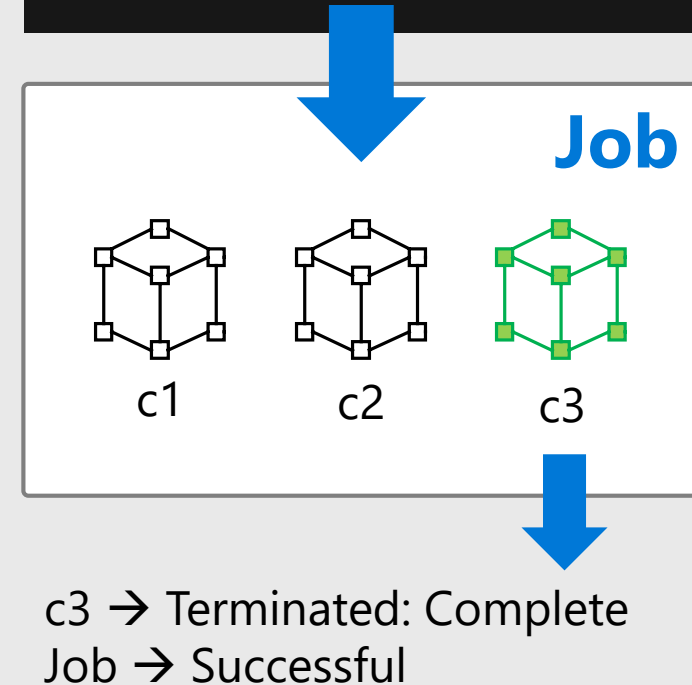
A *Job* creates one or more pods and ensures that a specified number of them successfully terminate

Jobs can run in **serial** or **parallel**

Control of number of failures, completions, restart policy and level of parallelism

*CronJobs* allows you to **schedule** *Jobs* to be run

```
kind: Job
spec:
  completions: 1
  parallelism: 3
```



Use Jobs for any workloads that run in batch or perform one off tasks

# Namespaces


## Multi tenancy

Most Kubernetes objects live inside a *Namespace*

Kubernetes starts with two *Namespaces*: **default** and **kube-system**

You can create *Namespaces* to **logically partition** a cluster, e.g. for dev/test or different customers

Nodes will be shared across *Namespaces*



When first learning Kubernetes use the default Namespace

# Role Based Access Control (RBAC)

## Regulating & governing access

RBAC controls user and system access to the API and Kubernetes resources

*Roles* define privileges as sets of **verbs** and API resources

*RoleBinding* grants *Roles* to *Users* and *ServiceAccounts*

Kubernetes has many authentication and authorization schemes. This complex area is deemed out of scope for this guide



RoleBinding



```
kind: Role
name: pod-reader
rules:
  resources: ["pods"]
  verbs: ["get", "watch", "list"]
```



Verb:

**get, delete, create, list, update, watch**

API:

**pods, secrets, services, jobs, nodes, ingresses**

RBAC is optional, but is becoming standard for any cluster

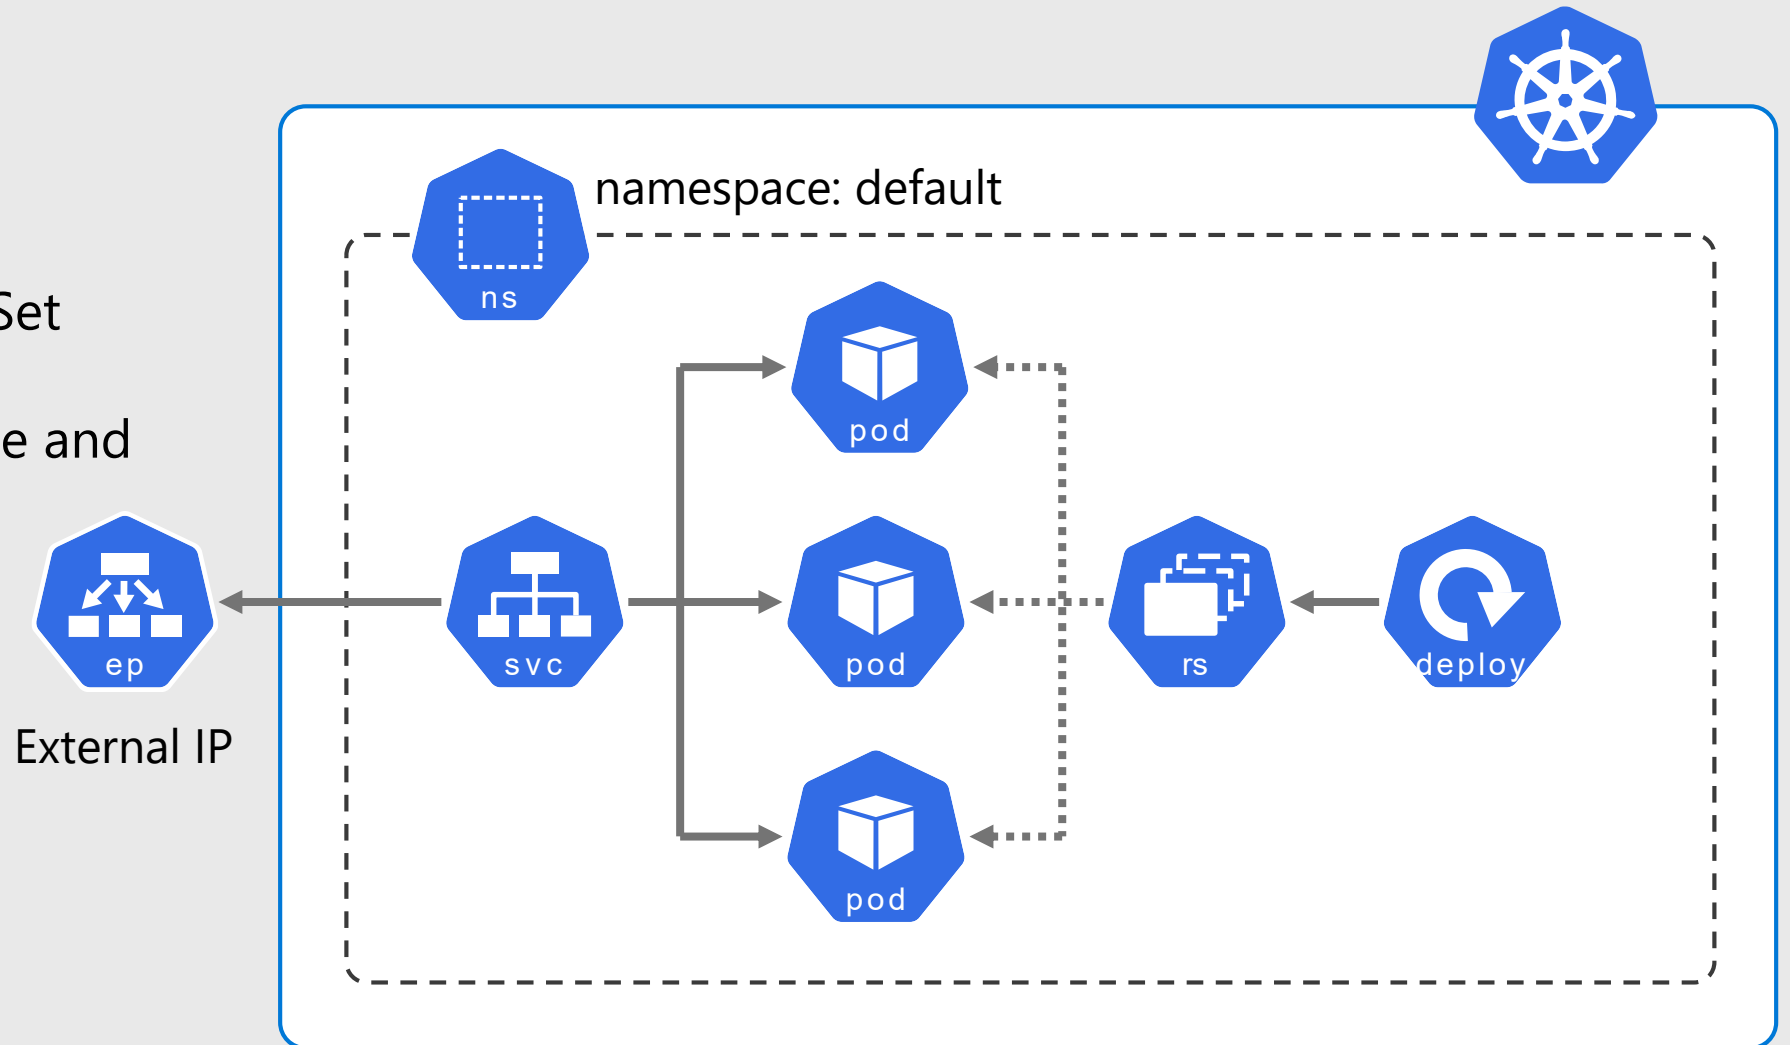
# Putting It All Together

## Sample Architecture – Simple App

### Basic web application

Running in 3 pods via a  
Deployment and ReplicaSet

With LoadBalancer service and  
external IP



# Using Kubernetes



# Command Line - kubectl

Main management & control  
interface to Kubernetes

Single executable binary

Can manage several clusters

Secure interface to Kubernetes API

Used for cluster ops and application  
deployment & management

```
ben@Azure:~$ kubectl get nodes
NAME                                STATUS    ROLES    AGE
aks-agentpool-25884925-0           Ready     agent     35d
aks-agentpool-25884925-1           Ready     agent     28d
aks-agentpool-25884925-2           Ready     agent     28d
ben@Azure:~$
ben@Azure:~$
ben@Azure:~$ kubectl run nginx --image nginx
deployment.apps/nginx created
ben@Azure:~$
ben@Azure:~$ kubectl get deploy
NAME            DESIRED    CURRENT    UP-TO-DATE    AVAILABLE
data-api        3          3          3             3
frontend        2          2          2             2
nginx           1          1          1             1
party-clippy    1          1          1             1
ben@Azure:~$
```



# Kubectl – Common Commands

get	delete	apply	run	describe
Display one or many resources	Delete resources	Create resources from YAML	Directly start and run pods	Get details of any resource
<pre>get nodes get pods get all get pod/myPod get nodes -w</pre>	<pre>delete deploy/mydeploy  delete -f myapp.yaml  delete -l app=foo</pre>	<pre>apply -f myapp.yaml</pre>	<pre>run nginx --image=nginx --replicas=3  run myapp --image=foo/img --port=8080</pre>	<pre>describe pod/pod138  describe svc/myservice  describe pod -l app=myapp</pre>



[kubernetes.io/docs/reference/kubectl/cheatsheet](https://kubernetes.io/docs/reference/kubectl/cheatsheet)

[kubernetes.io/docs/reference/kubectl/overview](https://kubernetes.io/docs/reference/kubectl/overview)

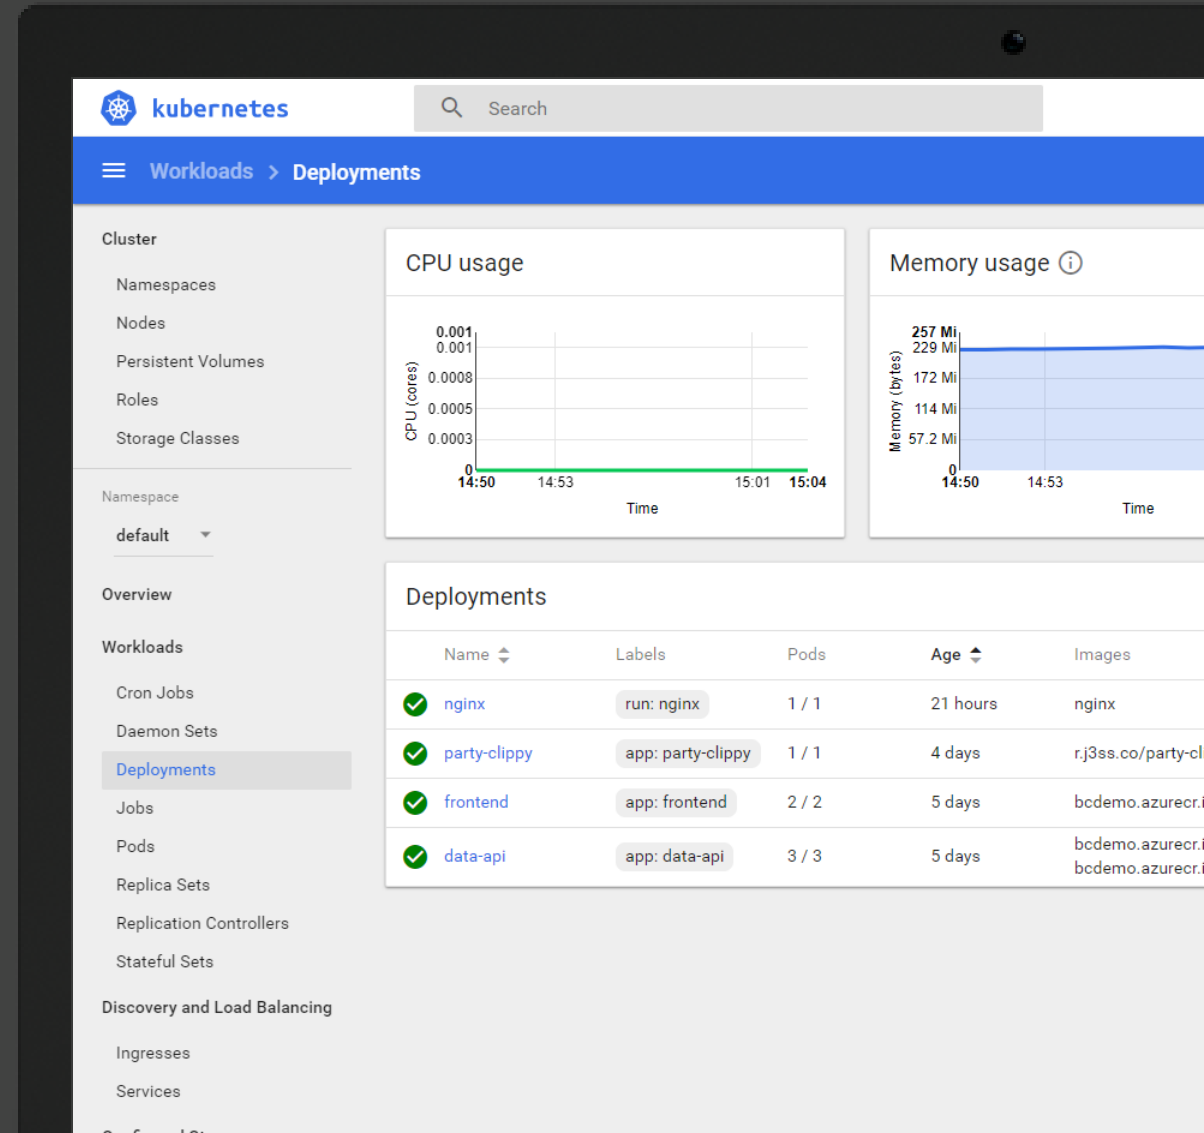
# Dashboard

Management web UI

Runs inside Kubernetes as a pod

Not secure, should never be exposed externally

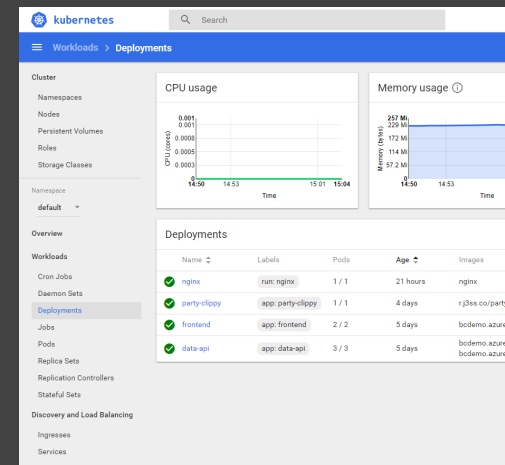

Optional, not required to run Kubernetes



# Accessing the Dashboard

- Open proxy to Kubernetes API server
- Run command: `kubectl proxy`
- Access this URL

<http://localhost:8001/api/v1/namespaces/kube-system/services/kubernetes-dashboard/proxy>



# Kubernetes Object Management

Three methods of managing Kubernetes



[kubernetes.io/docs/concepts/overview/object-management-kubectl](https://kubernetes.io/docs/concepts/overview/object-management-kubectl)

Management Technique	Operates On	Recommended Environment	Learning Curve	Infrastructure As Code
<b>Imperative commands</b>	Live objects	Dev projects	Lowest	No
<b>Imperative object configuration</b>	Individual files	Production use	Moderate	Limited
<b>Declarative object configuration</b>	Individual & multiple files	Production use	Highest	Yes

File based declarative object configuration  
is the most common approach used

# Introduction to the Declarative Model

YAML or JSON documents

Describe any Kubernetes object

Objects & properties map directly to the Kubernetes API

You can combine multiple objects into a single file (separate with ---)

Note. JSON is also supported, but YAML is recommended for readability

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: mydeploy
spec:
  replicas: 4
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
      - name: mycontainer
        image: bencuk/vuego-demoapp
        ports:
        - containerPort: 4000
```



This is a deployment object, called '**mydeploy**'



It will run 4 replicas of a pod matching the label **app=myapp**



Each pod will be labelled with **app=myapp** and



Runs a container from image **bencuk/vuego-demoapp**



Port 4000 will be exposed from the container

# Idempotent Updates & Desired State

Files can describe **desired state** of the objects you configuring

```
kubectl apply
```

Kubernetes applies updates in **idempotent** way, modifying objects only if needed

Idempotency: "The definition of the target state can be applied multiple times and if the system's state is unchanged, no changes are made to the system"

```
$ kubectl apply -f myconfig.yaml
```

```
deployment.apps/mydeploy created
```



```
$ nano myconfig.yaml
```



myapp.yaml


```
spec:  
  replicas: 4
```



```
spec:  
  replicas: 2
```

```
$ kubectl apply -f myconfig.yaml
```

```
deployment.apps/mydeploy configured
```



# Labels & Selectors



[kubernetes.io/docs/concepts/overview/working-with-objects/labels](https://kubernetes.io/docs/concepts/overview/working-with-objects/labels)

Kubernetes makes extensive use of labels and selectors

**Labels** are metadata on any object and are just `key:value` pairs of your choosing

**Selectors** are lookups that match one or more objects based on their **labels**

## Uses

- Which pods are in a service
- Which nodes to run a pod on
- Which pods are in a deployment
- Logically group & tag resources

```
kind: Deployment
...
metadata:
  labels:
    app: myapp
    type: frontend
```

---

```
kind: Service
spec:
  type: LoadBalancer
  selector:
    app: myapp
```

This is a deployment labels its pods with **app: myapp** & **type: frontend**

This service will look for any pods labelled with **app=myapp**

Note. There might be zero, one or many matches

# Configuring Basic Workloads





# Environmental Variables



[kubernetes.io/docs/tasks/inject-data-application/define-environment-variable-container/](https://kubernetes.io/docs/tasks/inject-data-application/define-environment-variable-container/)

Environmental variables are the standard way to **configure containers at runtime**

Containerized app consumes environmental variables in standard way

Key value pairs

Application/container specific

## Uses

- Application configuration
- Parameter passing

```
apiVersion: v1
kind: Pod
metadata:
  name: envar-demo
  labels:
    purpose: demonstrate-envvars
spec:
  containers:
  - name: envar-demo-container
    image: gcr.io/google-samples/node-hello:1.0
    env:
      - name: DEMO_GREETING
        value: "Hello from the environment"
      - name: DEMO_FAREWELL
        value: "Such a sweet sorrow"
```

# Secrets

Hold **sensitive information** such as passwords, certs and API keys

Don't place sensitive values as plain text in deployment files

Don't "bake" secrets into your container images

Can be mounted in pods as files or environmental variables

## Uses

- TLS certificates
- Application configuration
- Authentication with private registry



[kubernetes.io/docs/concepts/configuration/secret](https://kubernetes.io/docs/concepts/configuration/secret)

```
$ kubectl create secret generic my-secret  
--from-literal=connString='blahblah_this_is_secret'
```

secret/my-secret created

myapp.yaml

```
containers:  
  - name: my-web-server  
  env:  
    - name: DATABASE_CONNECTION_STRING  
      valueFrom:  
        secretKeyRef:  
          name: my-secret  
          key: connString
```



# ConfigMaps

Hold application **configuration data**

Key value pairs (like secrets), YAML or free format (e.g. XML, conf)

Pass to containers as env vars or mount as volume



[kubernetes.io/docs/tasks/configure-pod-container/configure-pod-configmap](https://kubernetes.io/docs/tasks/configure-pod-container/configure-pod-configmap)

```
$ kubectl create configmap my-config  
--from-file=/path/to/foobar.conf
```

configmap/my-config created

```
containers:  
- name: my-foobar-server  
  volumeMounts:  
  - name: config-vol  
    mountPath: /etc/config  
volumes:  
- name: config-vol  
  configMap:  
    name: my-config
```

myapp.yaml



Mount into container at given path

Reference to config map object

## Uses

- Application configuration

# Resource Management



[kubernetes.io/docs/concepts/configuration/manage-compute-resources-container](https://kubernetes.io/docs/concepts/configuration/manage-compute-resources-container)

Define compute (CPU & memory)  
**resource limits and requests** for  
containers

Allows Kubernetes to make **better  
scheduling placement decisions**

**Limits** are enforced, **requests** aren't

CPU resources are fractions of 1  
vCore

## Uses

- Efficient use of cluster resources
- Prevent rogue workloads starving the cluster
- Good practice

```
apiVersion: v1
kind: Pod
metadata:
  name: demo-app
spec:
  containers:
  - name: db
    image: mysql
    resources:
      requests:
        memory: "64M"
        cpu: "0.25"
      limits:
        memory: "512M"
        cpu: "2.0"
```

Certain capabilities such as auto scaling  
are dependant on setting resources

Specifying limits & requests is optional  
but **STRONGLY** recommended

# Liveness & Readiness Probes



[kubernetes.io/docs/tasks/configure-pod-container/configure-liveness-readiness-probes](https://kubernetes.io/docs/tasks/configure-pod-container/configure-liveness-readiness-probes)

Liveness probes tell Kubernetes your **container is "alive"**

Readiness probes tell Kubernetes your container is **accepting traffic**

Liveness probe failure can **restart the container**

HTTP, TCP and command checks

## Uses

- Maintain availability
- Restart/termination of unhealthy containers
- Efficient traffic routing

...

```
livenessProbe:
  httpGet:
    path: /status
    port: 8080
  initialDelaySeconds: 25
  periodSeconds: 10
  failureThreshold: 3
```

...

```
readinessProbe:
  exec:
    command: ["mysqladmin", "ping"]
  initialDelaySeconds: 30
  periodSeconds: 20
```

Specifying a liveness probe is optional but recommended

# Commands & Arguments



[kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/](https://kubernetes.io/docs/tasks/inject-data-application/define-command-argument-container/)

Using the **command** property pass a starting command to a container

Use **args** to pass arguments to the container, pass an array of strings

Note. These correspond to the Docker **Entrypoint** and **Cmd** parameters

```
kind: Pod
metadata:
  name: command-demo
spec:
  containers:
  - name: command-demo-container
    image: debian
    command: ["printenv"]
    args: ["HOSTNAME", "KUBERNETES_PORT"]
```

Uses
<ul style="list-style-type: none"><li>- Use base images to run utilities &amp; scripts</li><li>- Debugging &amp; trouble shooting</li><li>- Application configuration</li><li>- Parameter passing</li></ul>

Description	Docker field name	Kubernetes field name
The command run by the container	Entrypoint	command
The arguments passed to the command	Cmd	args

# Beyond The Basics



# Service Mesh

Facilitates **services to service** calls **inside** Kubernetes

## Features:

- Observability / tracing
- Error handling / retries / backoff
- Encryption / mTLS
- Routing / balancing

Typically runs as a proxy sidecar in all your pods

Adds complexity & overhead



Read more

<https://servicemesh.io/>




# Kubernetes Ecosystem – *Many* Projects


## Security / Policy



## DevOps



## Serverless



## Monitoring



## Network Plugins



## Load Balancing & Discovery



## Storage



[landscape.cncf.io](https://landscape.cncf.io)

# The Kubernetes API

It's Kind Of Important!



[kubernetes.io/docs/reference/  
#api-reference](https://kubernetes.io/docs/reference/#api-reference)

Every object in Kubernetes and any interactions with the cluster are shaped by the API and the API spec

- YAML manifests schema
- `kubectl` commands

Served by the API server running on master node(s)

Kubernetes version dependant

Aggregated - hosts multiple APIs at multiple versions

`kubectl proxy` – Create a local tunnel to the API server



# Additional Network Services



# Ingress

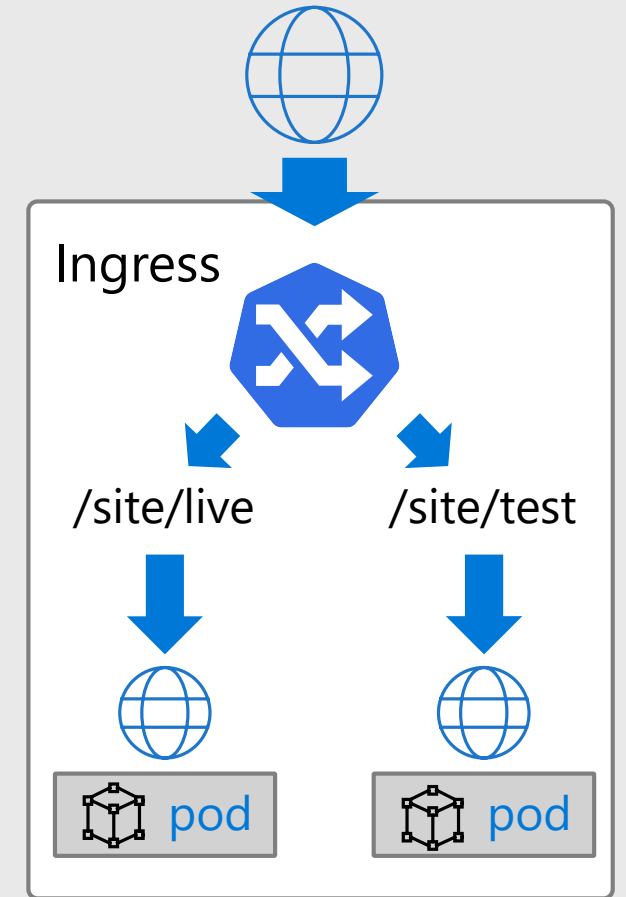
## External access for HTTP and HTTPS

An *Ingress* allows you to **route** HTTP/HTTPS traffic to **services** based on URL and/or domain host name

*Ingress* object is a **set of rules** picked up and implemented by the *Ingress Controller*

*Ingress Controller* has a **public IP** and *LoadBalancer* service, it routes traffic to internal *ClusterIP* services

Various controllers exist, NGINX is commonly used



Use an Ingress when you want to route HTTP(S) traffic into your workloads and pods

# External DNS

## Optional Addon – Auto configuration of public DNS

Allows for **dynamic configuration** of DNS records

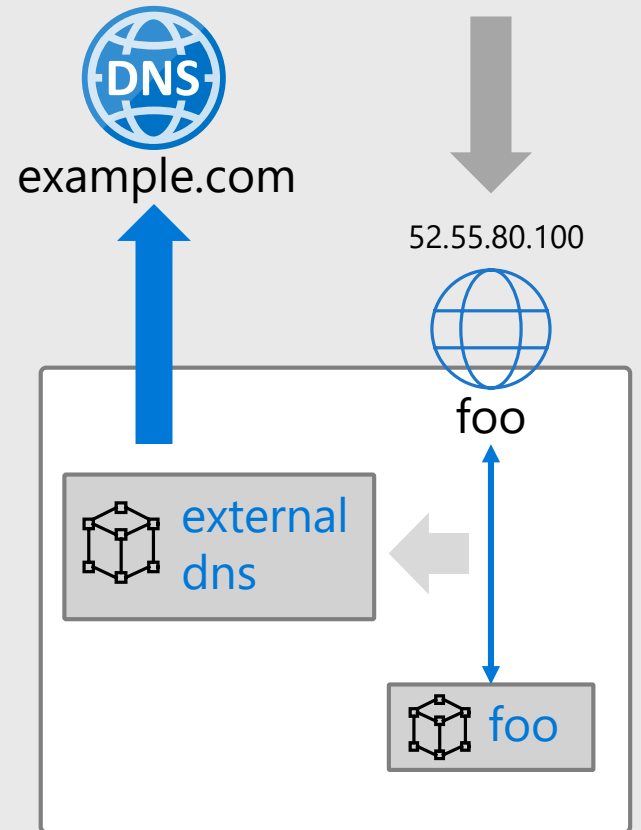
Seamlessly keep **public DNS** in sync with your Ingress and external services

Supports Azure DNS, AWS, CloudFlare, Google DNS etc

Incubation project:

[github.com/kubernetes-incubator/external-dns](https://github.com/kubernetes-incubator/external-dns)

A record: `foo.example.com`  
IP: `52.55.80.100`



Commonly used with an Ingress for host based external routing

# Cert Manager

## Optional Addon – Automate issuing of TLS certs

Ensures **certificates are valid** and up to date

Tightly coupled to *Ingress*, e.g. host rules

**Renew** certificates before expiry

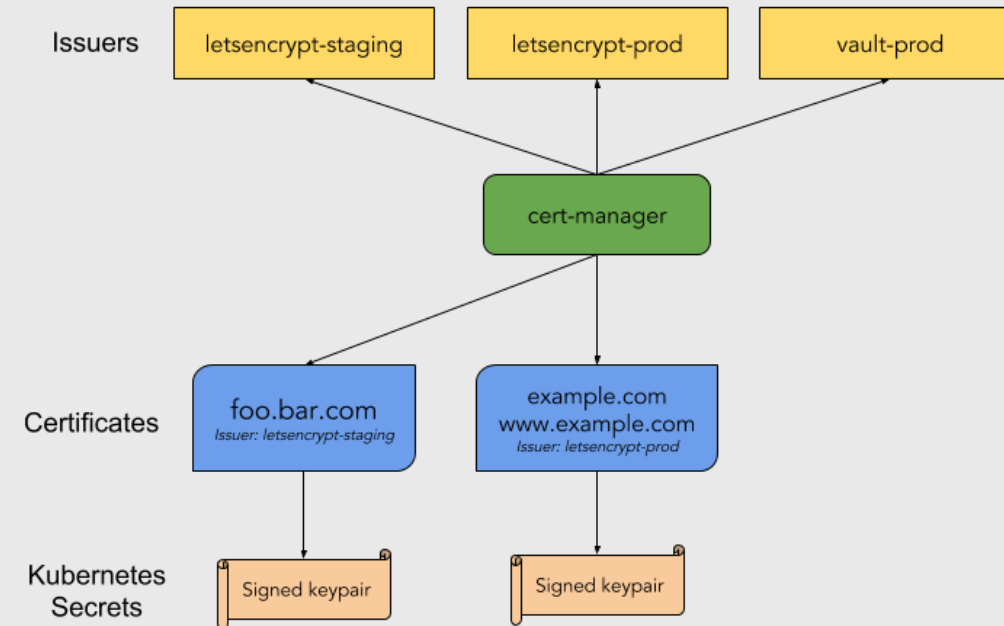
Uses ACME issuers, i.e. **Let's Encrypt**

Project:

[github.com/jetstack/cert-manager](https://github.com/jetstack/cert-manager)



[github.com/jetstack/cert-manager](https://github.com/jetstack/cert-manager)



Issue TLS certs for HTTPS access to services & Ingress

# Putting It All Together

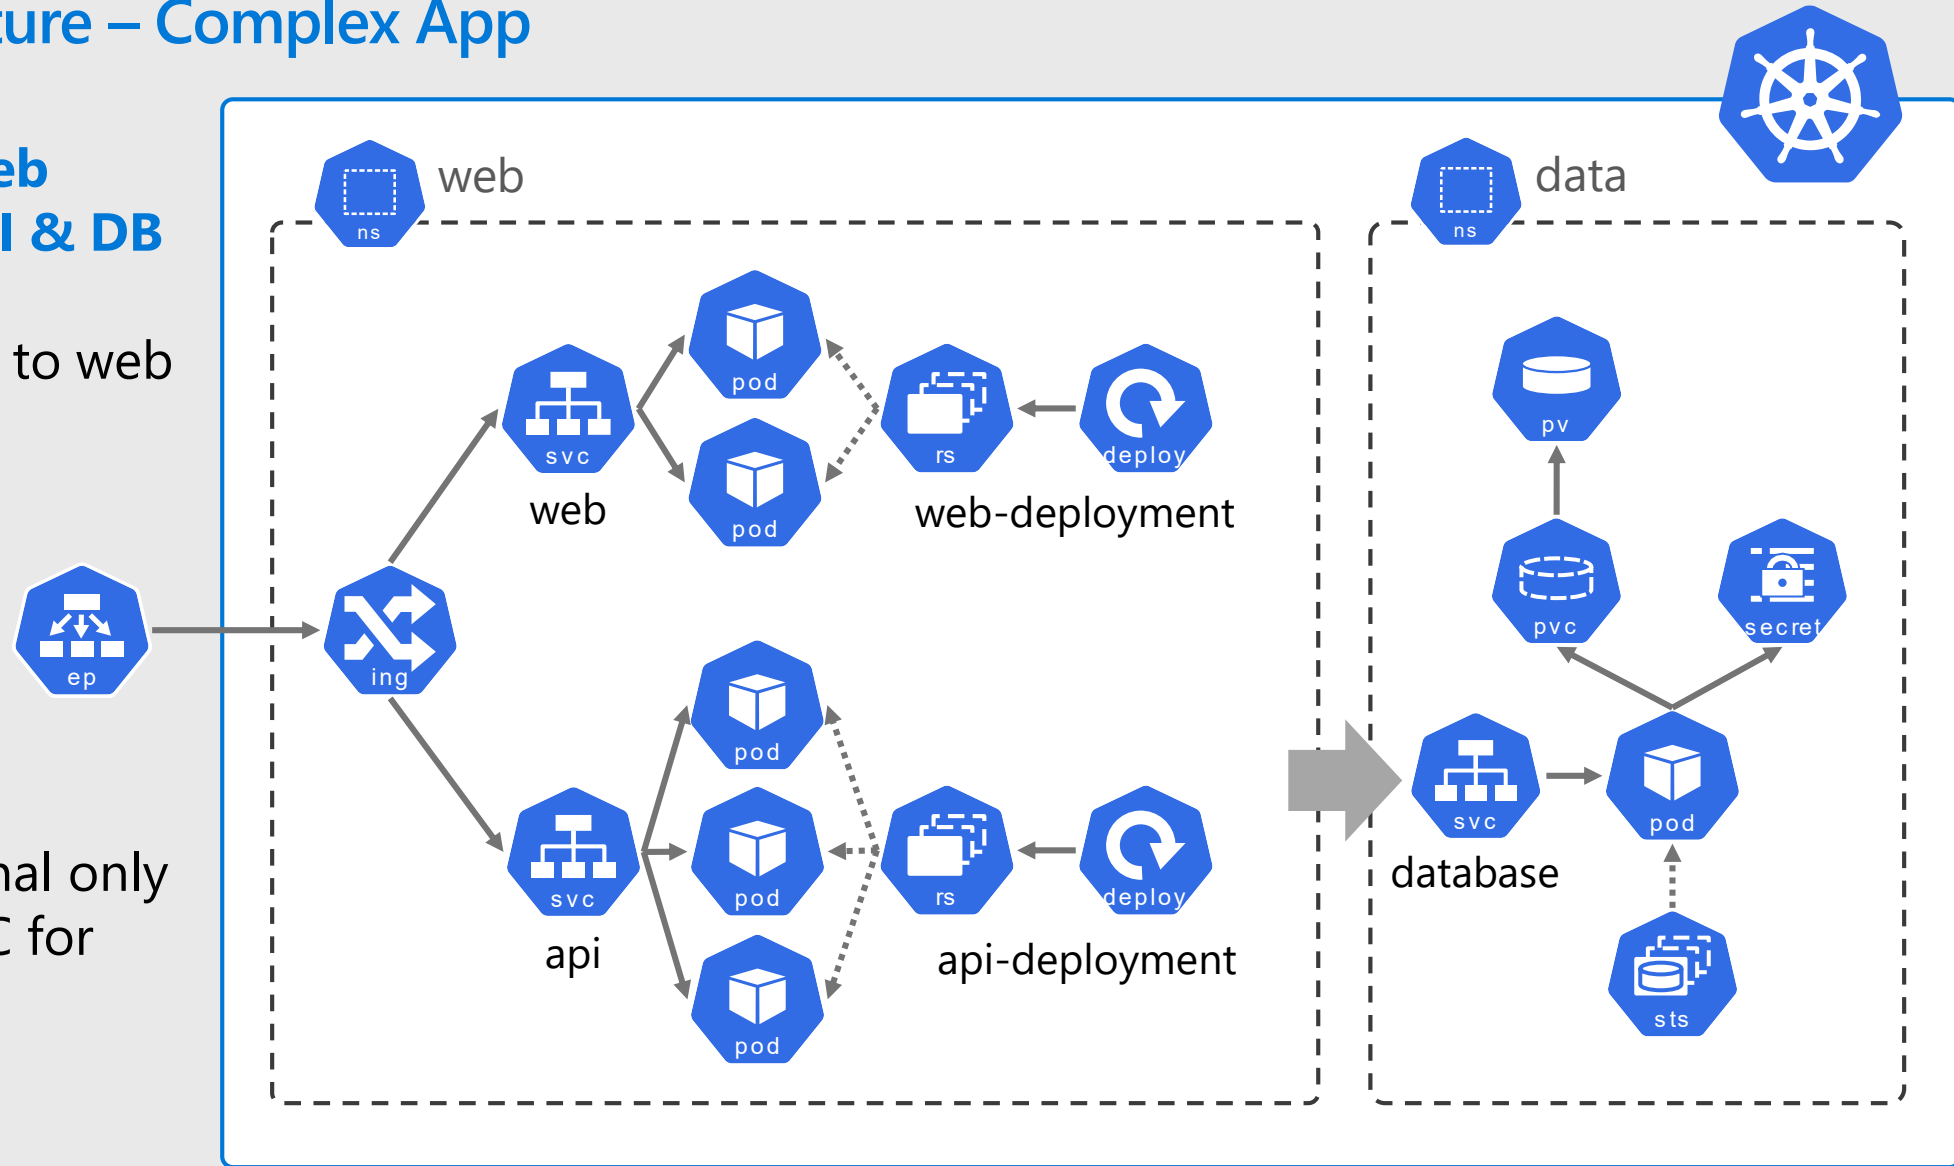
## Sample Architecture – Complex App

### Multi component web application, with API & DB

Ingress routing traffic to web and API pods

Ingress is exposed to the internet

Database running in StatefulSet with internal only service and using PVC for storage, plus secrets



# Debugging and Troubleshooting Workloads





# Describing Objects



kubernetes.io/docs/tasks/debug-application-cluster

Use `kubectl describe` to inspect status of any object in your cluster

Returns events and all properties & status details

Use label selectors to query multiple objects

## Uses

- View and understand the status of anything in your cluster
- Troubleshoot pending/failed workloads

```
$ kubectl describe deploy/data-api
$ kubectl describe pod/data-api-84fb56497b-6cgth
$ kubectl describe service/frontend
$ kubectl describe pod -l app=data-api
```

```
Limits:
  cpu:    100m
  memory: 256M
Requests:
  cpu:    100m
  memory: 256M
Liveness: http-get http://:4000/api/info delay=3s timeout=1s period=20s #success=1 #failure=3
Environment:
  MONGO_CONNSTR:      mongodb://mongodb-svc.default
  KUBERNETES_PORT_443_TCP_ADDR: bckube-d587b0d8.hcp.northeurope.azmk8s.io
  KUBERNETES_PORT:      tcp://bckube-d587b0d8.hcp.northeurope.azmk8s.io:443
  KUBERNETES_PORT_443_TCP: tcp://bckube-d587b0d8.hcp.northeurope.azmk8s.io:443
  KUBERNETES_SERVICE_HOST: bckube-d587b0d8.hcp.northeurope.azmk8s.io
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-hvcxv (ro)
Conditions:
  Type            Status
  Initialized      True
  Ready            False
  ContainersReady  True
  PodScheduled     True
Volumes:
  default-token-hvcxv:
    Type: Secret (a volume populated by a Secret)
    SecretName: default-token-hvcxv
    Optional: false
QoS Class:       Guaranteed
Node-Selectors:  <none>
Tolerations:     node.kubernetes.io/not-ready:NoExecute for 300s
                  node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type    Reason      Age   From                  Message
  ----    -
  Normal  Scheduled   22m   default-scheduler     Successfully assigned default/data-api-84fb56497b-6cgth to aks-agentpool-25884925-3
  Normal  Pulling     22m   kubelet, aks-agentpool-25884925-3   pulling image "bcdemo.azurecr.io/smlr/data-api"
  Normal  Pulled      22m   kubelet, aks-agentpool-25884925-3   Successfully pulled image "bcdemo.azurecr.io/smlr/data-api"
  Normal  Created     22m   kubelet, aks-agentpool-25884925-3   Created container
  Normal  Started     22m   kubelet, aks-agentpool-25884925-3   Started container
  Normal  Pulling     21m   kubelet, aks-agentpool-25884925-3   pulling image "bcdemo.azurecr.io/smlr/data-api:stable"
  Normal  Pulled      21m   kubelet, aks-agentpool-25884925-3   Successfully pulled image "bcdemo.azurecr.io/smlr/data-api:stable"
  Normal  Created     21m   kubelet, aks-agentpool-25884925-3   Created container
```

# Container Logs

Access stdout & stderr output  
from pods with:

`kubectl logs`

Get output from a `deployment` or  
`pod` or `single container`

Follow logs with `-f`

Uses
<ul style="list-style-type: none"><li>- View any errors output from containers</li><li>- See what your workloads are doing</li></ul>



```
$ kubectl logs deploy/data-api
```

```
Found 3 pods, using pod/data-api-84fb56497b-6cgth
```

```
> smilr-data-api@3.2.0 start /home/app
```

```
> node server.js
```

```
### Node environment mode is 'production'
```

```
### Connection attempt 1 to MongoDB server mongodb-  
svc.default
```

```
### Yay! Connected to MongoDB server
```

```
### Server listening on 4000
```

# Get Shell Access

Create **interactive shell** into running containers

Use **kubectl exec**

Use **-it** switch and sh or bash for shell

Can also run single non-interactive command

## Uses

- Low level debugging
- Interactive troubleshooting



[kubernetes.io/docs/tasks/debug-application-cluster/get-shell-running-container](https://kubernetes.io/docs/tasks/debug-application-cluster/get-shell-running-container)

```
$ kubectl exec frontend-58b84f7d7-fb6rg -- ps -ef
```

PID	USER	TIME	COMMAND
1	root	0:00	npm
22	root	0:00	node server.js
36	root	0:00	ps -ef

```
$ kubectl exec -it frontend-58b84f7d7-fb6rg -- bash
```

```
bash-4.4# ls -la
node_modules      package.json      server.js
```

```
bash-4.4# uname
Linux
```

# Advanced Pod Configuration



# Deeper Dive on Manifests

Manifests for *Deployments*, *StatefulSets* and *DaemonSets* have a similar pattern & structure

The **spec** part contains **replicas** and a **selector** and also a **template** for the objects it will replicate

The **template** will contain another **spec**, typically a *Pod spec*

A *Pod spec* contains one or more **containers**

?? So what do all these labels mean ??

```
kind: Deployment
apiVersion: apps/v1
metadata:
  name: my-deployment
  labels:
    cheese: cheddar
spec:
  replicas: 1
  selector:
    matchLabels:
      thing: my-app
  template:
    metadata:
      labels:
        thing: my-app
        cake: chocolate
    spec:
      containers:
      - name: my-container
        image: nginx
```

Deployment name will also be used as name prefix for the ReplicaSet and Pods

Labels for the deployment are optional

The selector of the spec **MUST** match one label in the template

Required field but mostly not important

# Init Containers



[kubernetes.io/docs/concepts/workloads/pods/init-containers](https://kubernetes.io/docs/concepts/workloads/pods/init-containers)

Init Containers are optional special containers that **run only once** when a pod is started

Init containers run to completion (**terminate**)

Main containers in a pod will **not start** until all Init Containers have run

## Uses

- Application configuration
- Bootstrapping apps
- Running utility & start-up scripts
- Data injection

...

initContainers:

- name: init-mysql  
image: bencuk/mysqlldb  
command: ["../scripts/checkDB"]

...

initContainers:

- name: init-demodata  
image: bcdemo.azurecr.io/smilr/data-api  
command: ['sh', '-c', 'cd demoData && node demodb.js']  
env:
  - name: MONGO\_CONNSTR  
value: mongodb://mongodb-svc.default
  - name: WIPE\_DB  
value: "true"

# Node Selector



[kubernetes.io/docs/concepts/configuration/assign-pod-node/#nodeselector](https://kubernetes.io/docs/concepts/configuration/assign-pod-node/#nodeselector)

A simple **constraint** to which Nodes are eligible to run a Pod

Key value pairs of labels, to be matched against Node's labels

Not a 'hard rule', other Pods that have no nodeSelector **can still land on the node**

```
kind: Pod
metadata:
  name: machineLearning
spec:
  containers:
  - name: trainModel
    image: ml-image:latest
  nodeSelector:
    hardware: gpu
```

*Pod*



Can be scheduled on

```
Name:      aks-nodepool1-18655374-vmss000000
Roles:     agent
Labels:    agentpool=nodepool1
           beta.kubernetes.io/arch=amd64
           hardware=gpu
```

*Node*

## Uses

- Assign workloads requiring special hardware or resources, e.g. GPU
- Physical partitioning of cluster
- Separating noisy Pods

# Affinity and Taints



[kubernetes.io/docs/concepts/configuration/taint-and-toleration/](https://kubernetes.io/docs/concepts/configuration/taint-and-toleration/)

[kubernetes.io/docs/concepts/configuration/assign-pod-node/#affinity-and-anti-affinity](https://kubernetes.io/docs/concepts/configuration/assign-pod-node/#affinity-and-anti-affinity)

**Affinity** and **anti-affinity** provide advanced ways to control Pod placement

- 'Hard' rules and 'soft' rules
- Weighting and expressions

**Taints** and **Tolerations**. Taints are applied to *Nodes* and will not accept *Pods* that don't have a matching **Toleration**

## Uses

- Assign workloads requiring special hardware or resources, e.g. GPU
- Physical partitioning of cluster
- Separating noisy Pods

```
$ kubectl taint nodes myNode001 team=team1:NoSchedule
node/myNode001 tainted
```

```
kind: Pod
spec:
  tolerations:
  - key: team
    operator: Exists
    value: "team1"
    effect: NoSchedule
```

Unlike taints & nodeSelectors, affinity rules can also apply to **pods**. i.e. do or don't schedule these pods together

```
kind: Pod
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
      preferredDuringSchedulingIgnoredDuringExecution:
    podAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
      preferredDuringSchedulingIgnoredDuringExecution:
```



# Sidecars

Pods **co-locate multiple containers** together, sharing network and storage


**Sidecar** is a pattern where additional containers provide enhancing/optional capabilities

Bind containers together to form a single cohesive unit of service

## Uses

- Decompose architecture
- Build services incrementally
- Bolt on features

```
kind: Pod
metadata:
  name: monitored-webapp
spec:
  containers:
  - name: webserver
    image: nginx
  - name: log-monitor
    image: my-log-monitor
    args: ["--log-dir", "/var/logs/nginx"]
```



# Scaling



# Manually Scaling

Scale stateless workloads by controlling the number of **replicas** of a Deployment

Be careful scaling **StatefulSets**, unless the workload/application is "cluster aware"


**DaemonSets** don't require scaling

Uses
<ul style="list-style-type: none"><li>- Horizontally scale</li><li>- Distribute work around cluster</li><li>- Remove single points of failure</li></ul>



```
$ kubectl scale deploy/myApp --replicas=5  
  
deployment.extensions/myApp scaled
```

OR

```
kind: Deployment  
apiVersion: apps/v1  
metadata:  
  name: myApp  
spec:  
  replicas: 5 
```

# Horizontal Pod Autoscaler (HPA)

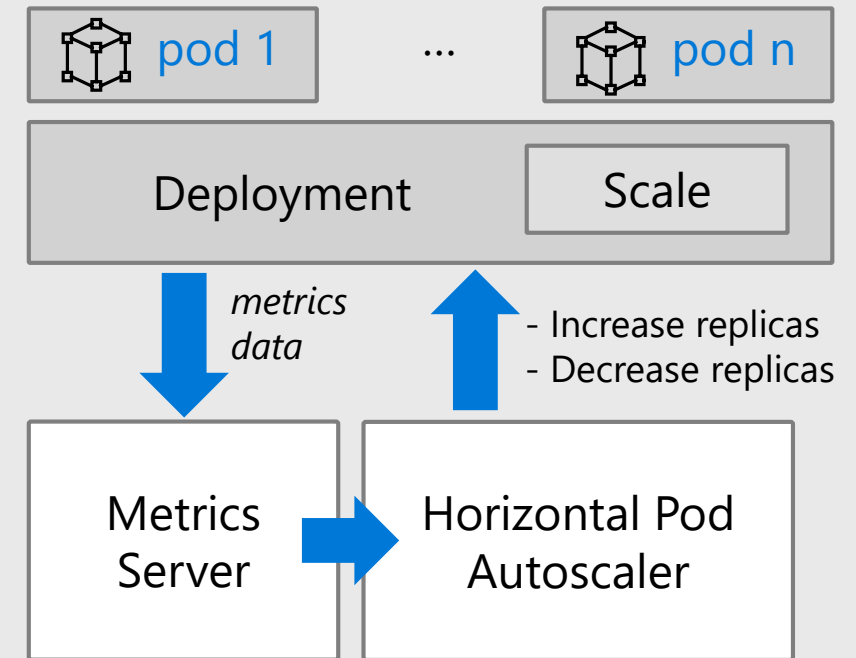
## Automatically scale stateless workloads

Use to horizontally **scale stateless** *Pods* in *ReplicaSet/Deployment*

Rules define desired *Pod* replicas based on **observed metrics**

Takes metrics from the metrics API fed from the *Metrics Server* (Kubernetes 1.8+)

Supports extension via custom metrics



**Dynamically scale stateless workloads across available nodes**


# Cluster Autoscaler (CA)

## Automatically scale cluster resources

Adjusts the size of the **Kubernetes cluster**, adding & removing *Nodes* when


- Pods are **pending** state due to insufficient resources (**scale out**)
- Nodes have been underutilized for a period of time (**scale in**)

Tightly coupled to the cloud and environment hosting the cluster & Nodes


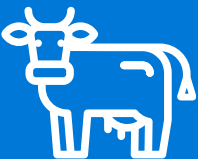


Scale cluster wide by adding/removing Nodes

# DevOps



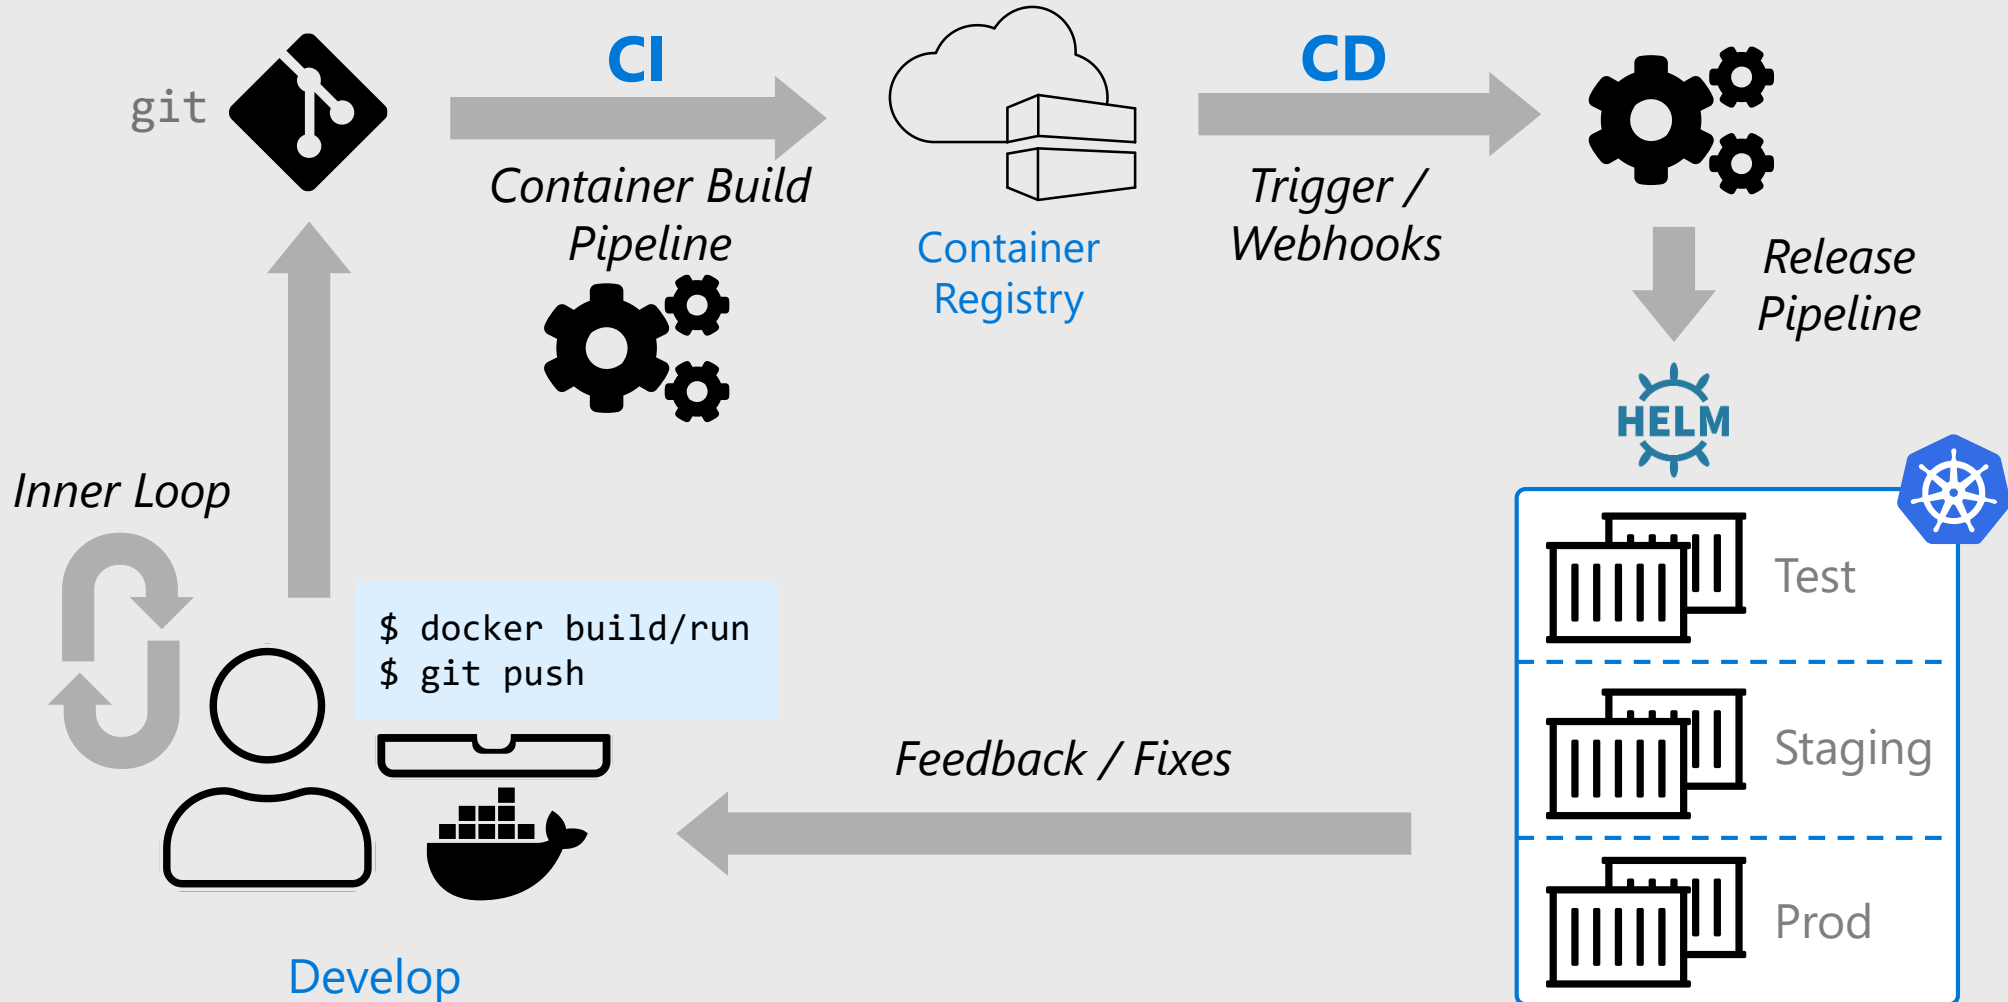
# Virtual Machines

<b>Pets</b> Legacy Infrastructure 	<b>Cattle</b> Modern Infrastructure 
<ul style="list-style-type: none"><li>Pets are given names like cutecat.petshop.com</li><li>They are unique, lovingly raised &amp; cared for</li><li>When they get ill, you nurse them back to health</li></ul>	<ul style="list-style-type: none"><li>Cattle are given numbers like 10200713.cattlerancher.com</li><li>They are almost identical to other cattle</li><li>When they get ill, you replace them and get another</li></ul>
Infrastructure is a <b>permanent</b> fixture in the data centre	Infrastructure is stateless, <b>ephemeral</b> , and <b>transient</b>
Infrastructure takes days to create, are <b>served weekly, maintained</b> for years, and requires migration project to move	Infrastructure is instantiated, modified, destroyed and recreated in minutes <b>from scratch</b> using <b>automated</b> scripts
<b>Infrastructure is modified</b> & patched in place and generally requires <b>specialist config management tools</b>	Infrastructure is <b>IMMUTABLE</b> Modify via <b>source control</b> & <b>automation pipelines</b> to rebuild and redeploy
Infrastructure requires <b>several different teams</b> to coordinate and provision the full environment	Infrastructure is <b>self-service</b> with the ability to provision computing, network and storage services with a single click
Infrastructure is <b>static</b> , requiring excess capacity to be dormant for use during peak periods of demands	Infrastructure is <b>elastic and scales automatically</b> , expanding and contracting on-demand to service peak usage periods

# Containers

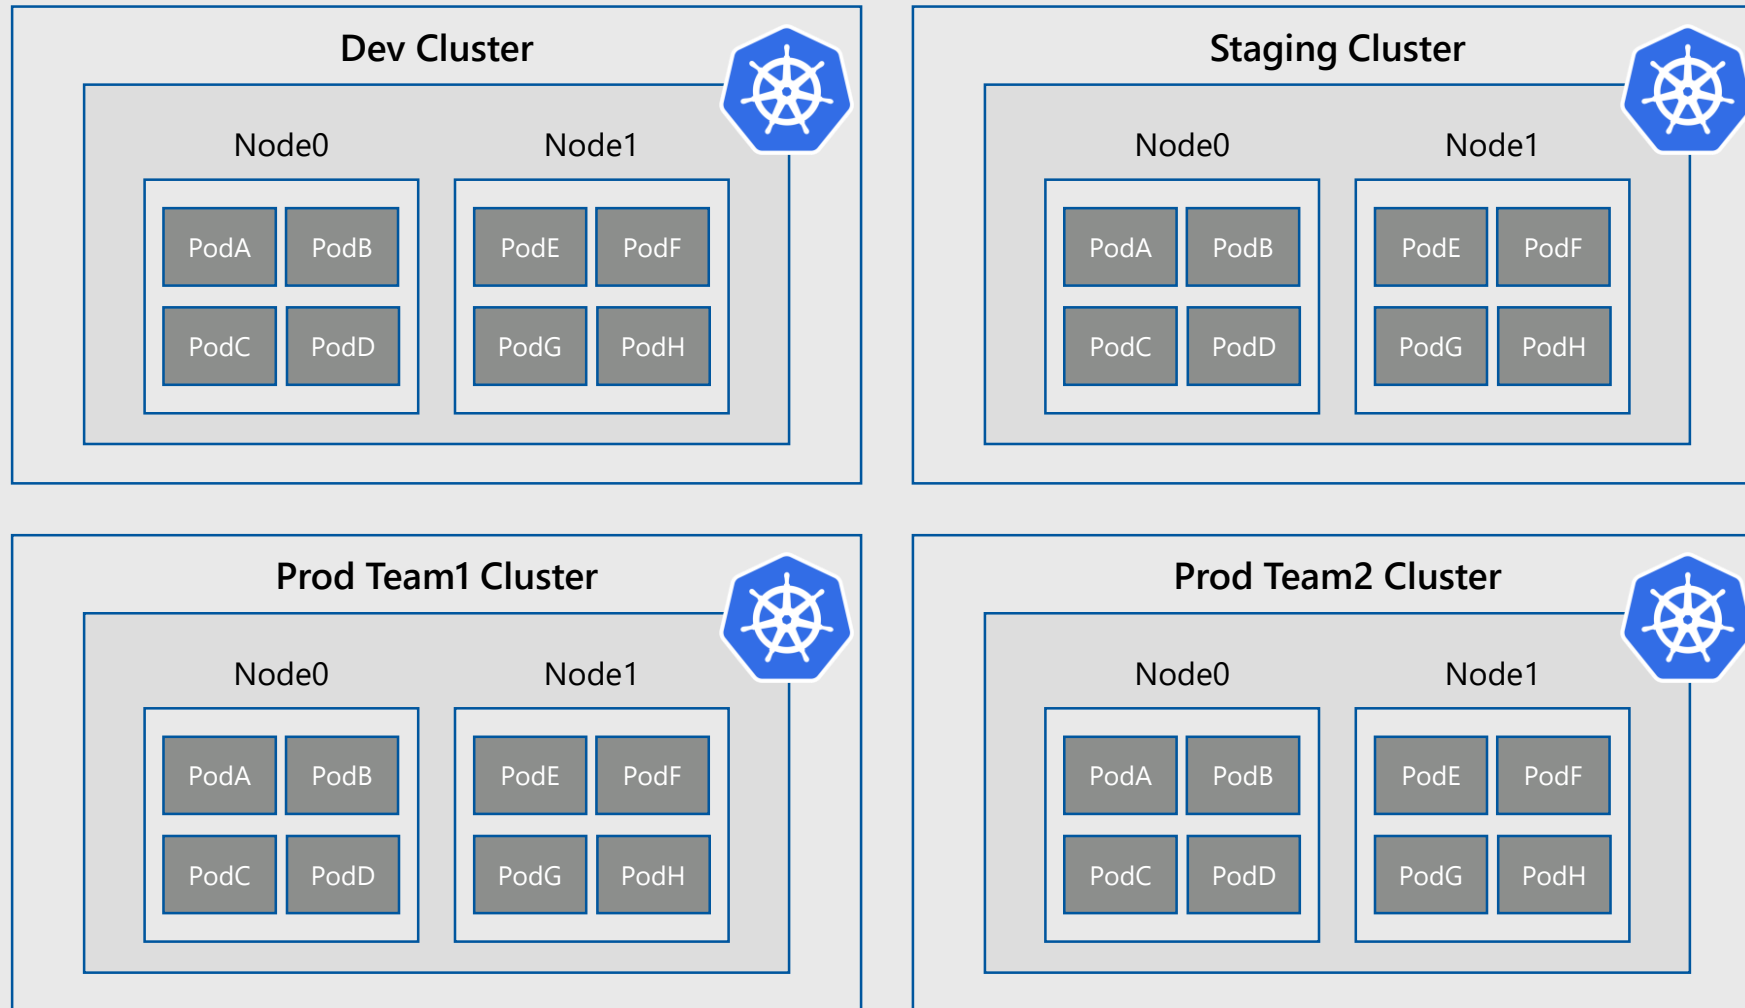
# DevOps Containers Lifecycle Loop

## Complete CI/CD Pipeline Loop



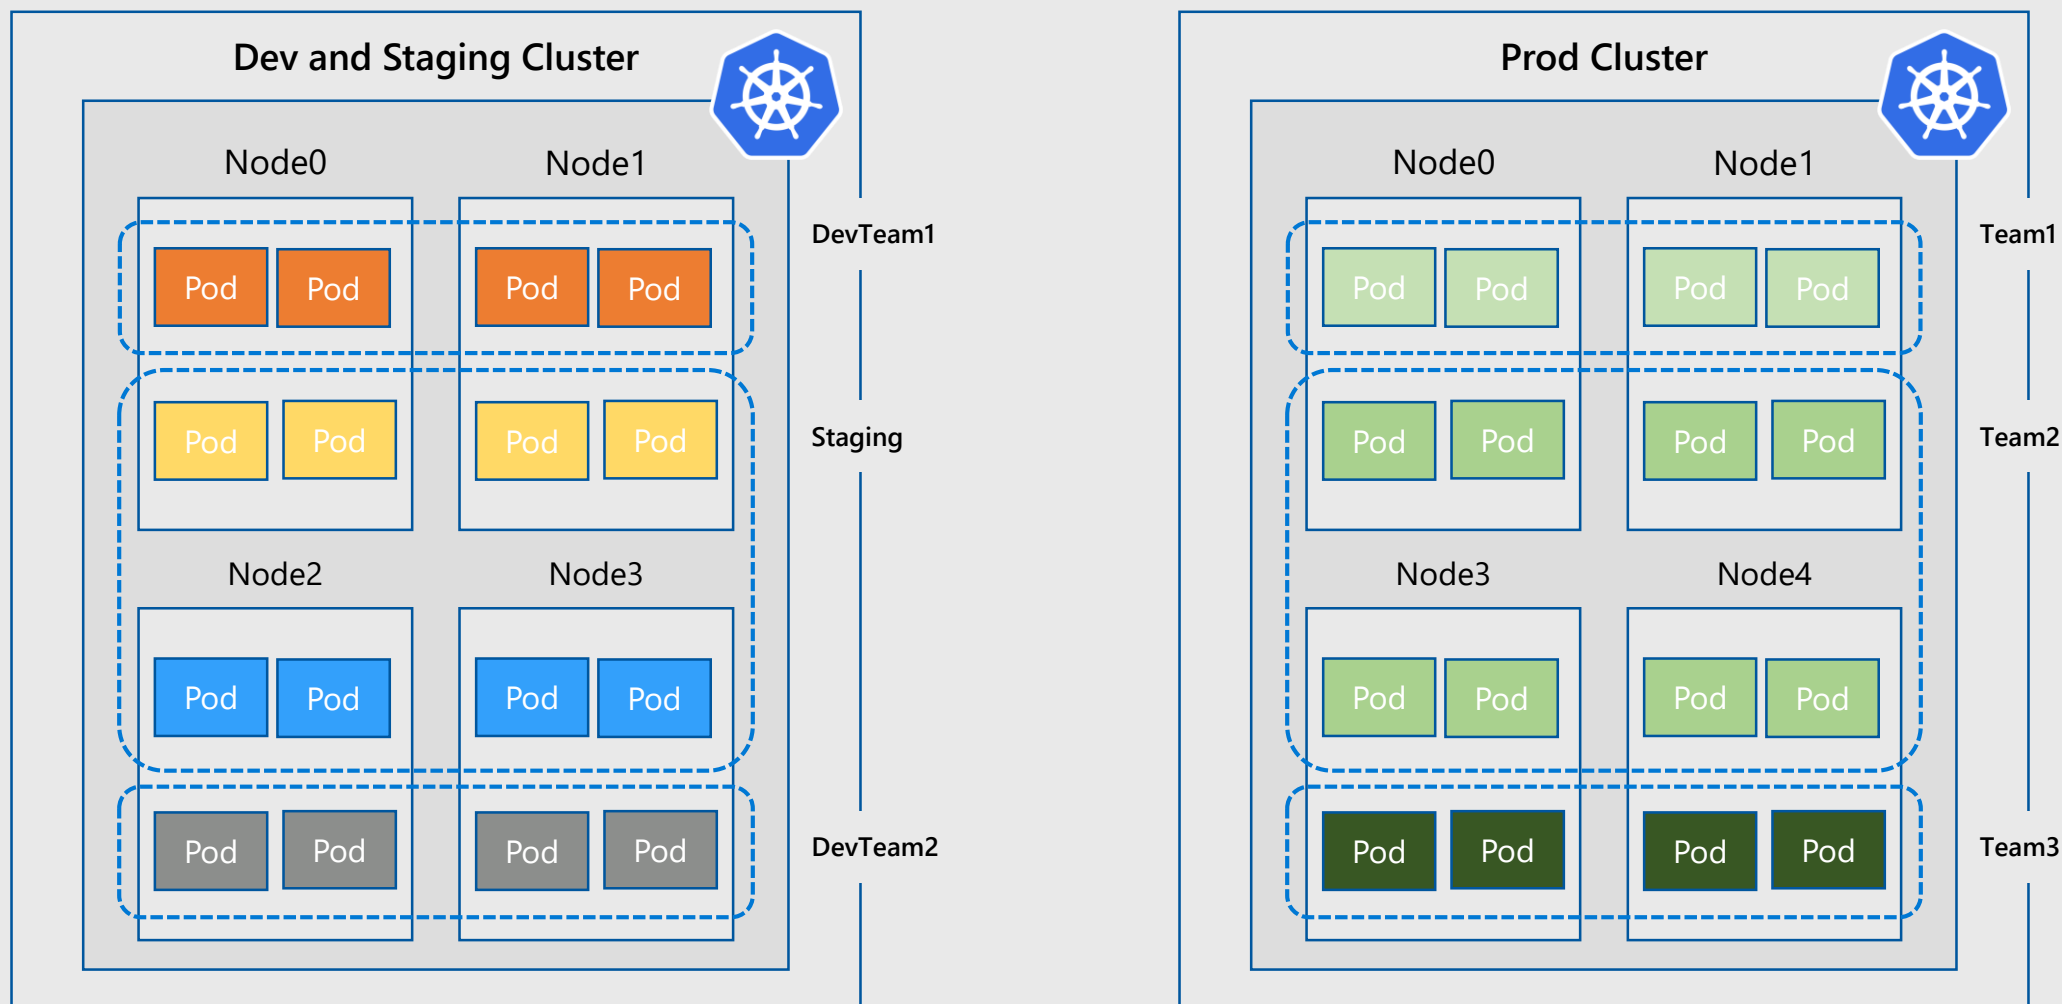


# Cluster Isolation Patterns: Physical Isolation



Simple but can waste resources, management overhead

# Cluster Isolation Patterns: Logical Isolation



Requires planning & some Kubernetes expertise

# Helm

## Package Manager for Kubernetes

Helm simplifies **deployment** into Kubernetes using *charts*

A chart consists of one or more Kubernetes YAML **templates** + supporting files

Helm charts support dynamic **parameters & functions** important for automated pipeline deployments

240+ charts exist for standard software/tools/packages  
[github.com/helm/charts](https://github.com/helm/charts)



<https://helm.sh>

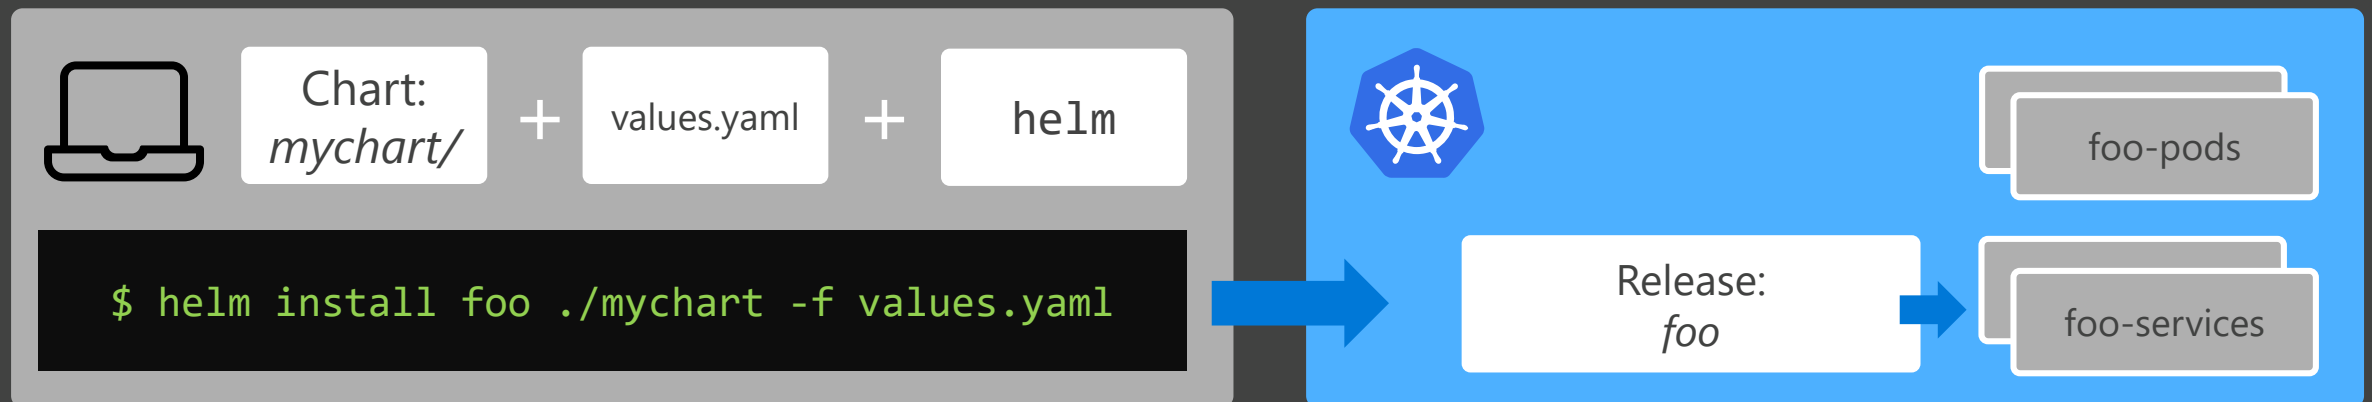


- Use Helm to install software/apps in your cluster
- Create Helm charts for your own apps, for CI/CD releases

# Helm – The Basics

[docs.helm.sh/glossary](https://docs.helm.sh/glossary)

<b>helm</b>	Client tool to manage and work with Helm
Chart	Package of Kubernetes resources in template form
Template	Kubernetes YAML with directives in Go template language format, e.g. {{ blah }}
Release	When installing a chart into Kubernetes it becomes a release
Values	Used at install time to customise the release, either from CLI or file
Dependency	A chart can require other external charts, Helm will automatically pull/update

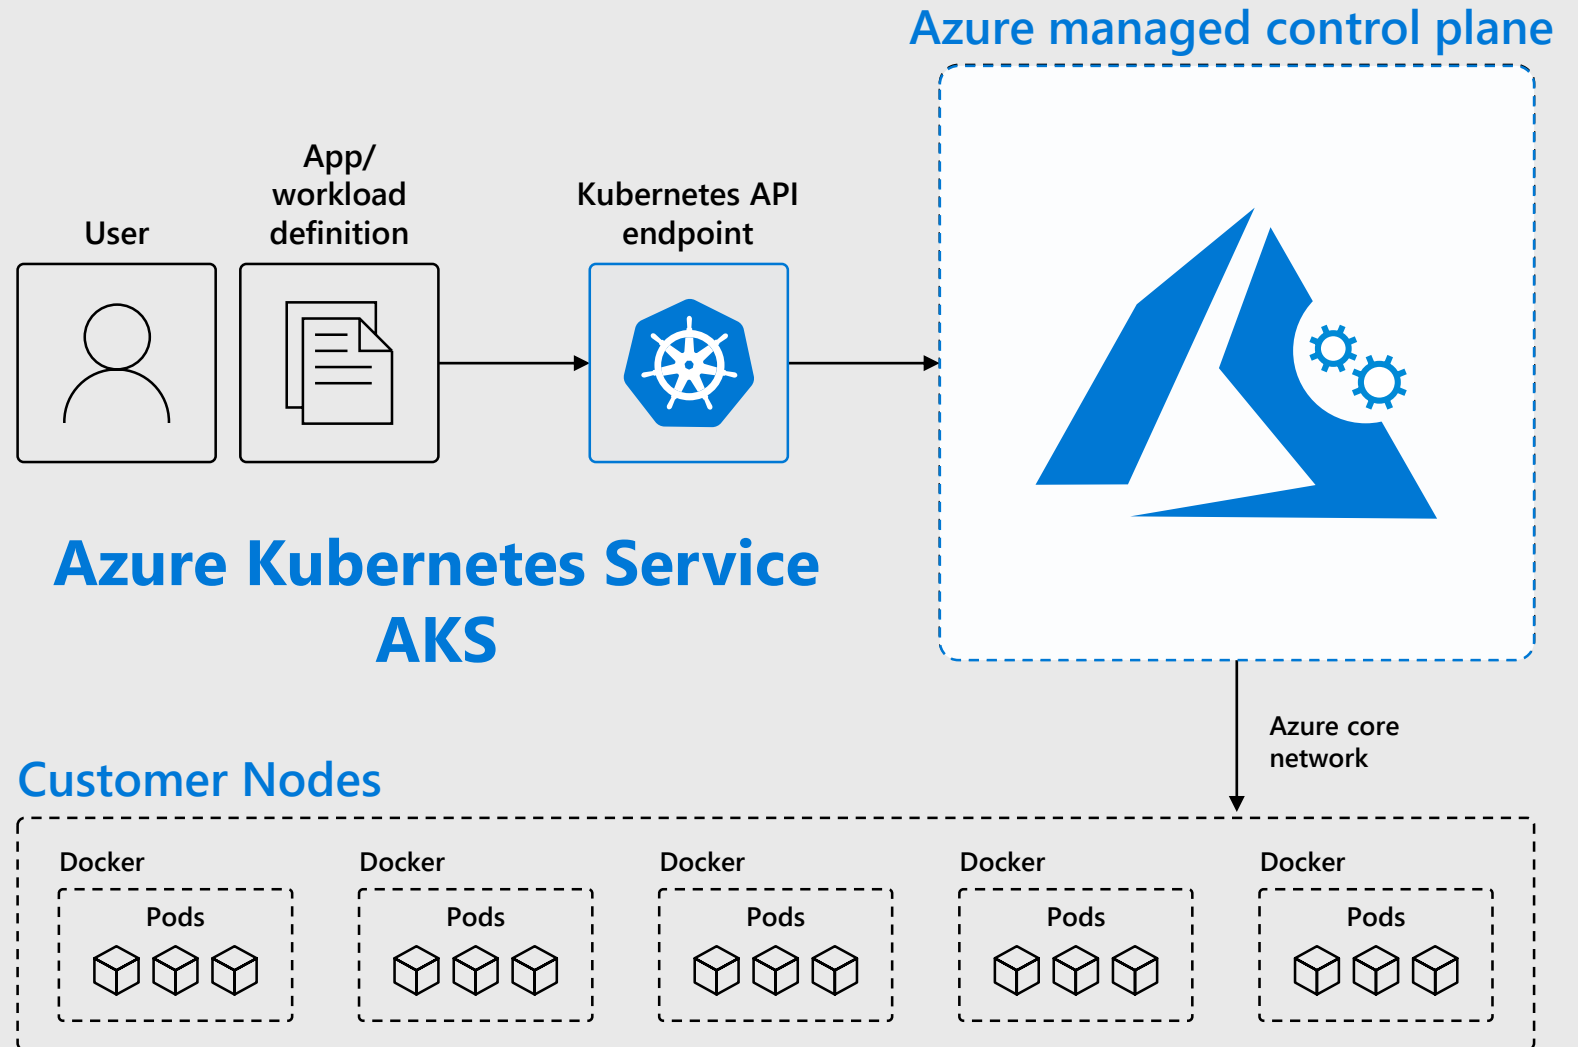


# Azure Kubernetes Service



# Managed Kubernetes on Azure – AKS

- Automated upgrades, patches
- High reliability, availability
- Easy, secure cluster scaling
- Self-healing
- API server monitoring
- At no charge



# AKS Virtual Nodes

Elastically provision compute capacity in seconds

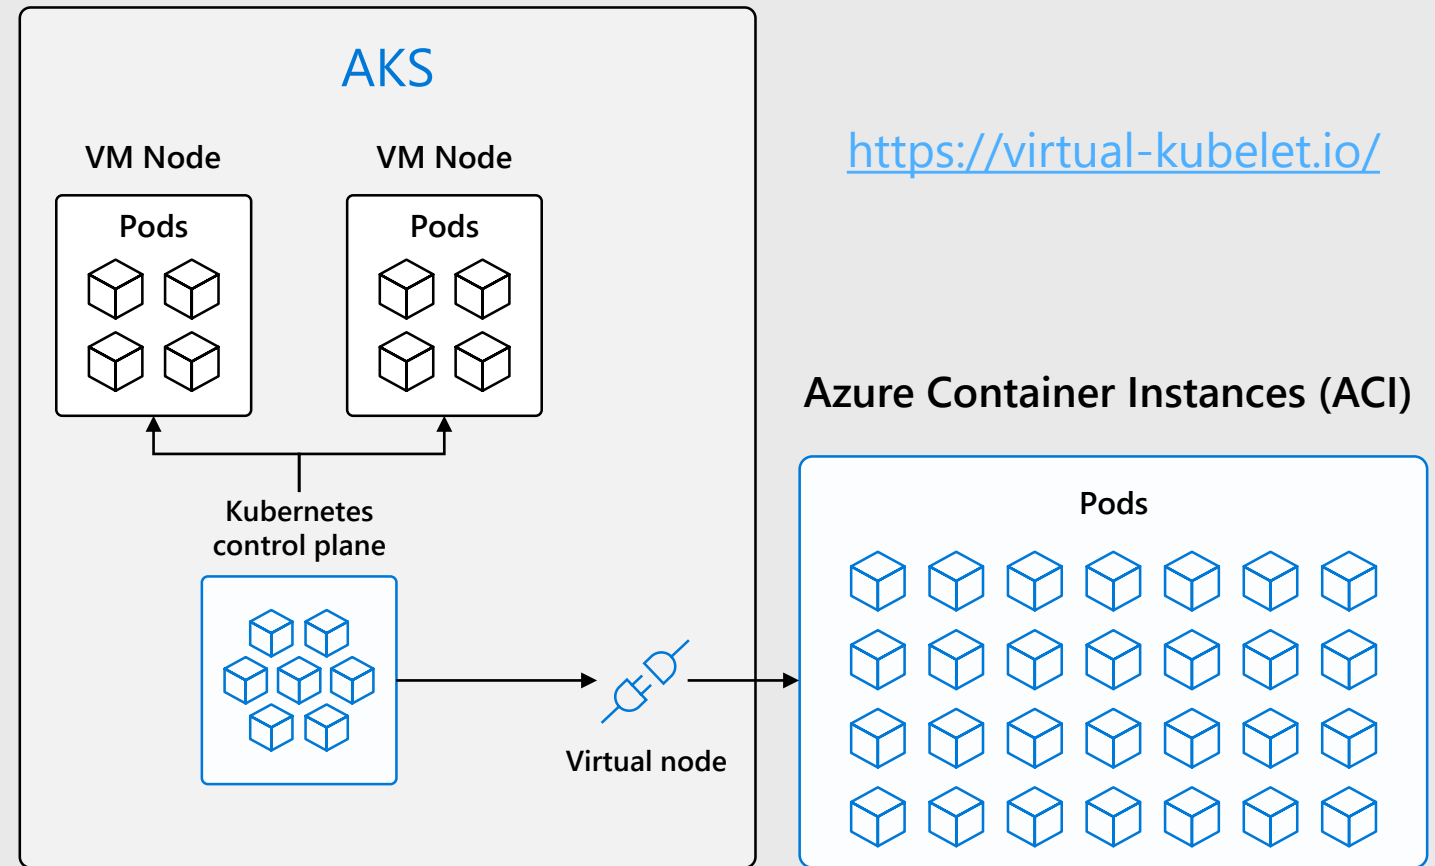
No infrastructure to manage

Builds on top of Azure Container Instances (ACI)

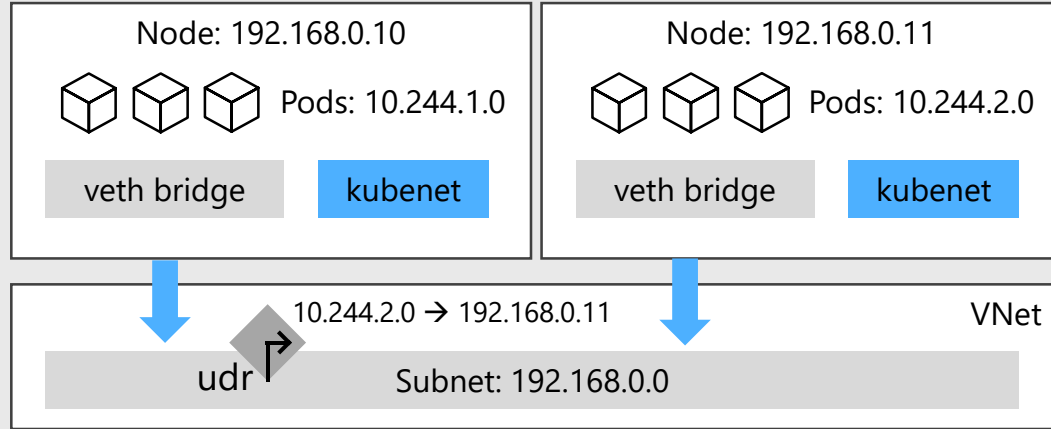
Uses Virtual Kubelet

Scenarios:

- Bursting
- CI/CD
- "Serverless Kubernetes"



# AKS Networking Models



Basic


Uses **kubenet** network plugin

Pods are not in a VNet or same subnet as Nodes

Pods live behind bridge, and UDR is used

Simple but has limitations & performance impact

No need to allocate/reserve IPs for Pods



Advanced

Uses **Azure CNI** network plugin

Both Nodes and Pods on same subnet + VNet

Works with peering, service endpoints, ExpressRoute

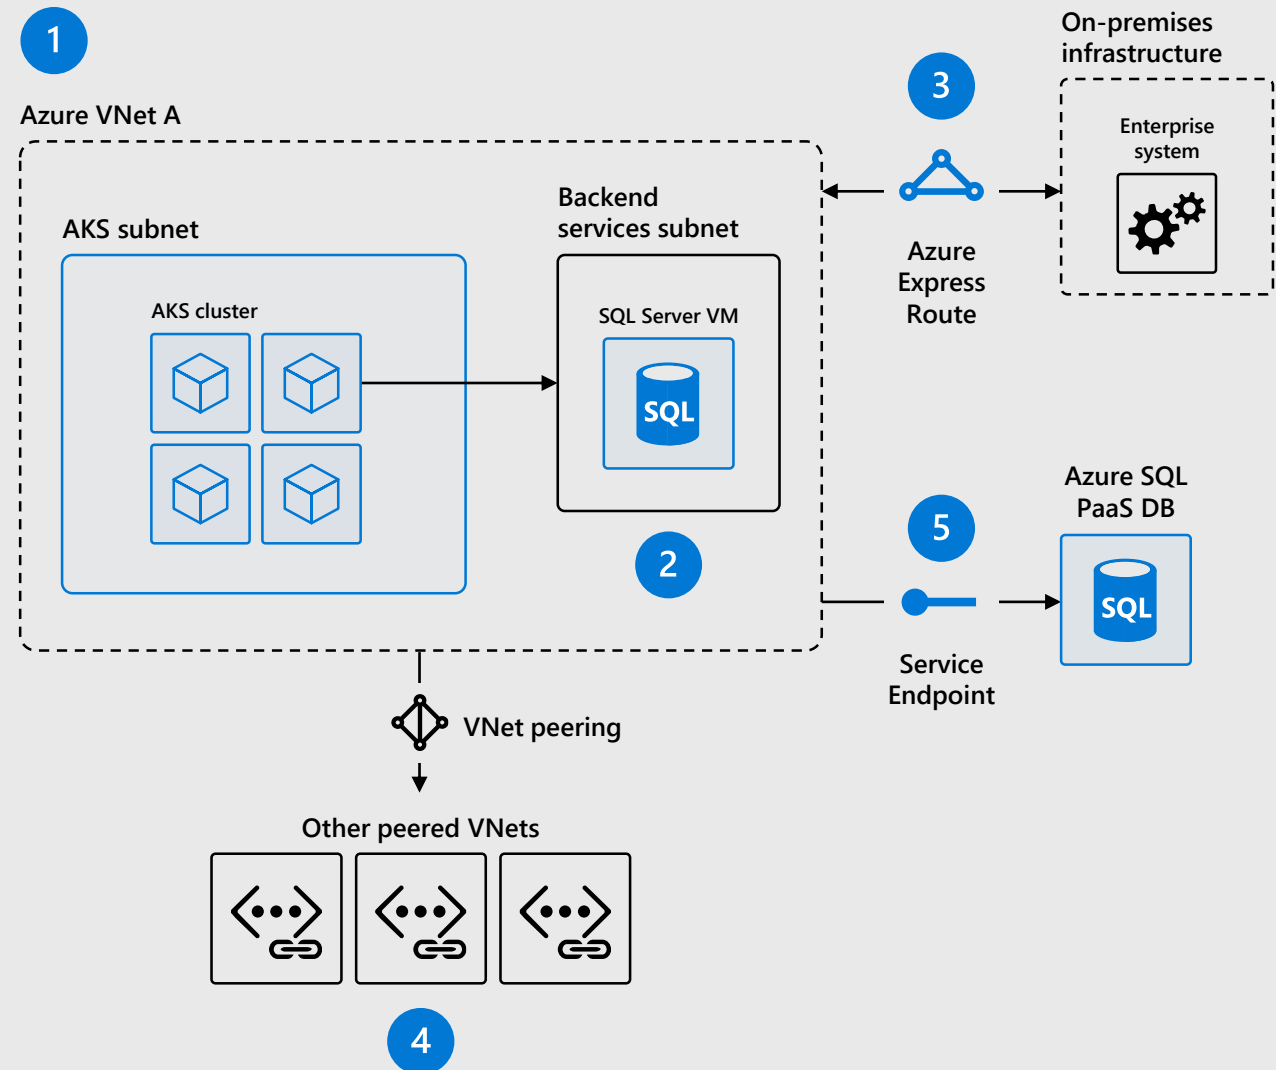
Better performance

Each pod requires an IP allocated to it




# AKS Advanced Networking

1. Uses Azure subnet for both your pods and cluster VMs
2. Allows for connectivity to existing Azure IaaS/VMs in the same VNet
3. Use ExpressRoute to connect to on-premises infrastructure
4. Use VNet peering to connect to other VNets
5. Connect AKS cluster securely and privately to Azure PaaS using VNet endpoints



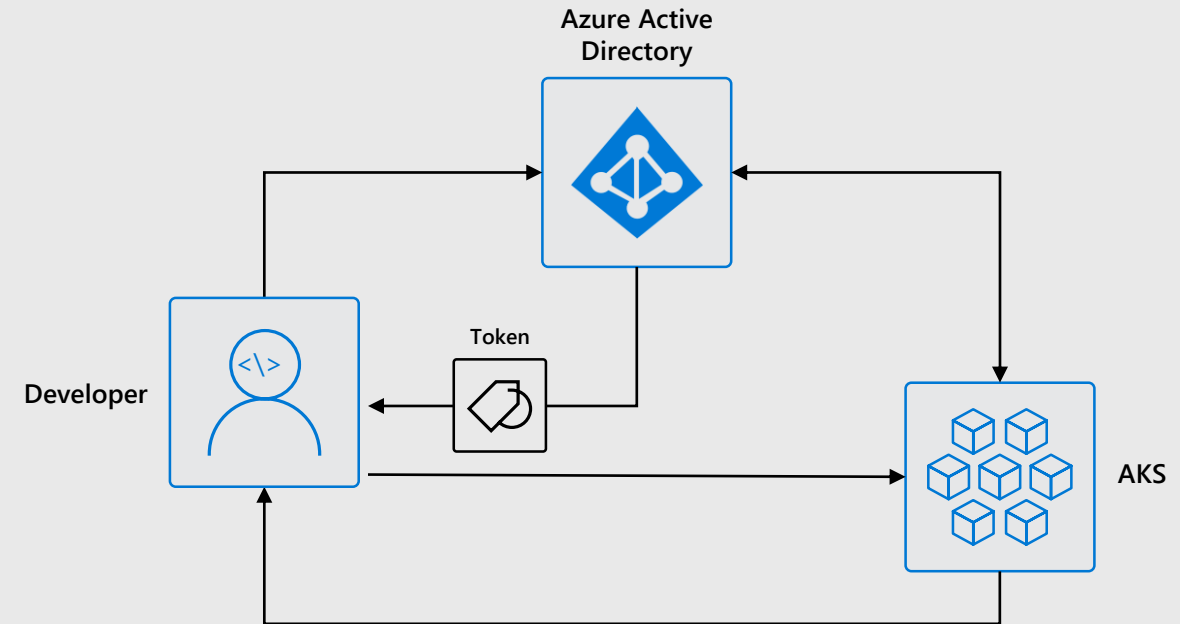
# Azure Active Directory

 [docs.microsoft.com/azure/aks/aad-integration](https://docs.microsoft.com/azure/aks/aad-integration)

Use Azure Active Directory to authenticate AKS users

Assign roles & permissions in Kubernetes RBAC, based on AAD users & groups

Roles in AAD define who has access to the cluster



Azure Kubernetes Service Cluster Admin Role



Azure Kubernetes Service Cluster User Role

Note. Can only be enabled at cluster creation time


# 'HTTP application routing' Add-On

 [docs.microsoft.com/azure/aks/  
http-application-routing](https://docs.microsoft.com/azure/aks/http-application-routing)

Simplifies Ingress and DNS for external access to HTTP services in cluster

Creates:

- NGINX Ingress controller (in cluster)
- External DNS helper (in cluster)
- aksapp.io DNS Zone (in Azure)



```
--enable-addons  
http_application_routing
```

The HTTP application routing add-on is designed to let you quickly create an ingress controller and access your applications. This add-on is not recommended for production use



# 'Monitoring' Add-On

## Azure Monitor – Container Insights

### Visualization

Visualize overall health and performance from clusters to containers with drill downs and filters

### Insights

Provide insights with multi-cluster health roll up view

### Monitor & Analyze

Monitor and analyze Kubernetes and container deployment performance, events, health, and logs

### Response

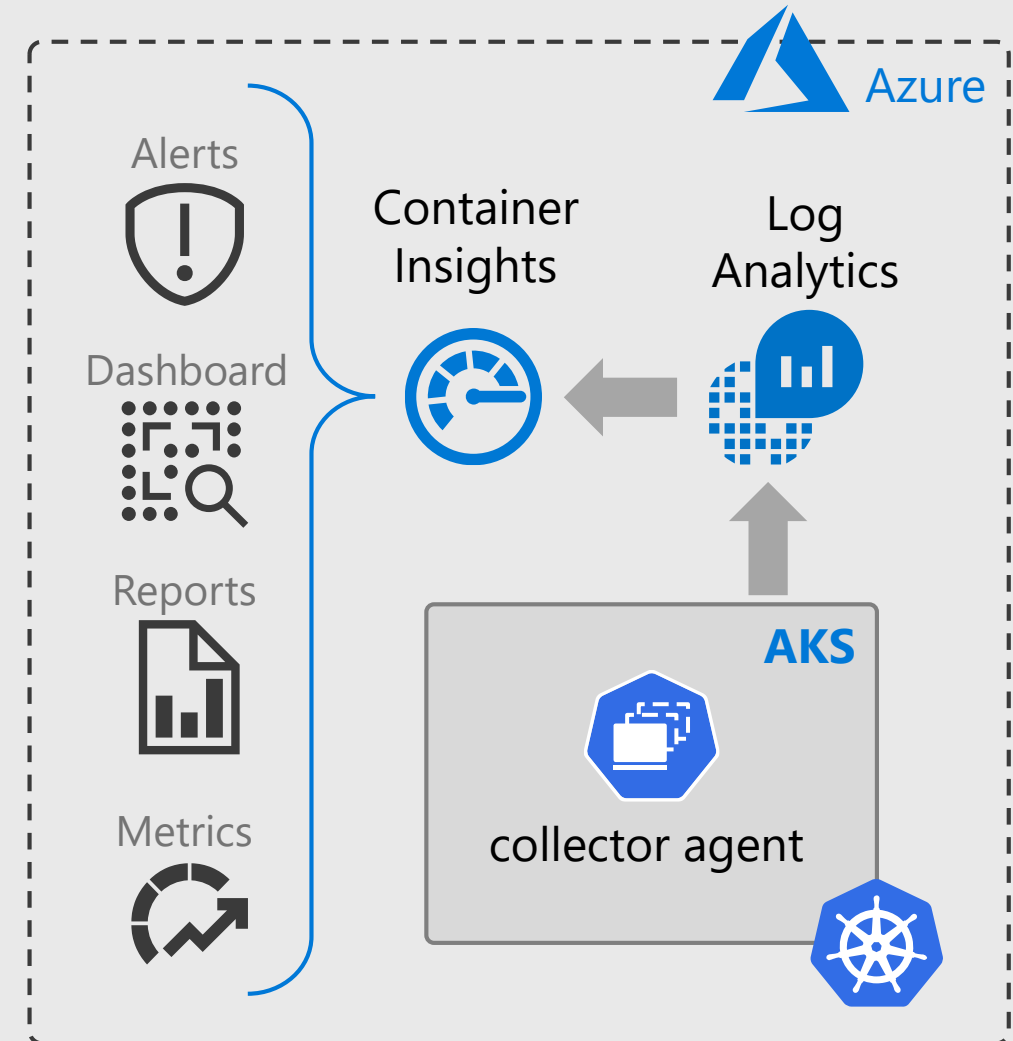
Native alerting with integration to issue managements and ITSM tools

### Observability

Observe live container logs on container deployment status



[docs.microsoft.com/azure/azure-monitor/insights/container-insights-overview](https://docs.microsoft.com/azure/azure-monitor/insights/container-insights-overview)



```
--enable-addons monitoring
--workspace-resource-id
{my-azure-monitor-logs-ws}
```

# Cluster Auto Scaler

## Autoscaling with Azure Scale Sets

Enables the standard Kubernetes Cluster Autoscaler (CA)

Uses VM Scale Sets (Preview)

While in preview - hidden behind CLI extension

```
--enable-vmss  
--enable-cluster-autoscaler  
--min-count 1  
--max-count 8
```




[docs.microsoft.com/en-us/azure/aks/cluster-autoscaler](https://docs.microsoft.com/en-us/azure/aks/cluster-autoscaler)

PREVIEW

Azure managed control plane



Node Pool



# Azure Dev Spaces

## Simplifying Development in Kubernetes



[docs.microsoft.com/en-us/azure/dev-spaces/](https://docs.microsoft.com/en-us/azure/dev-spaces/)

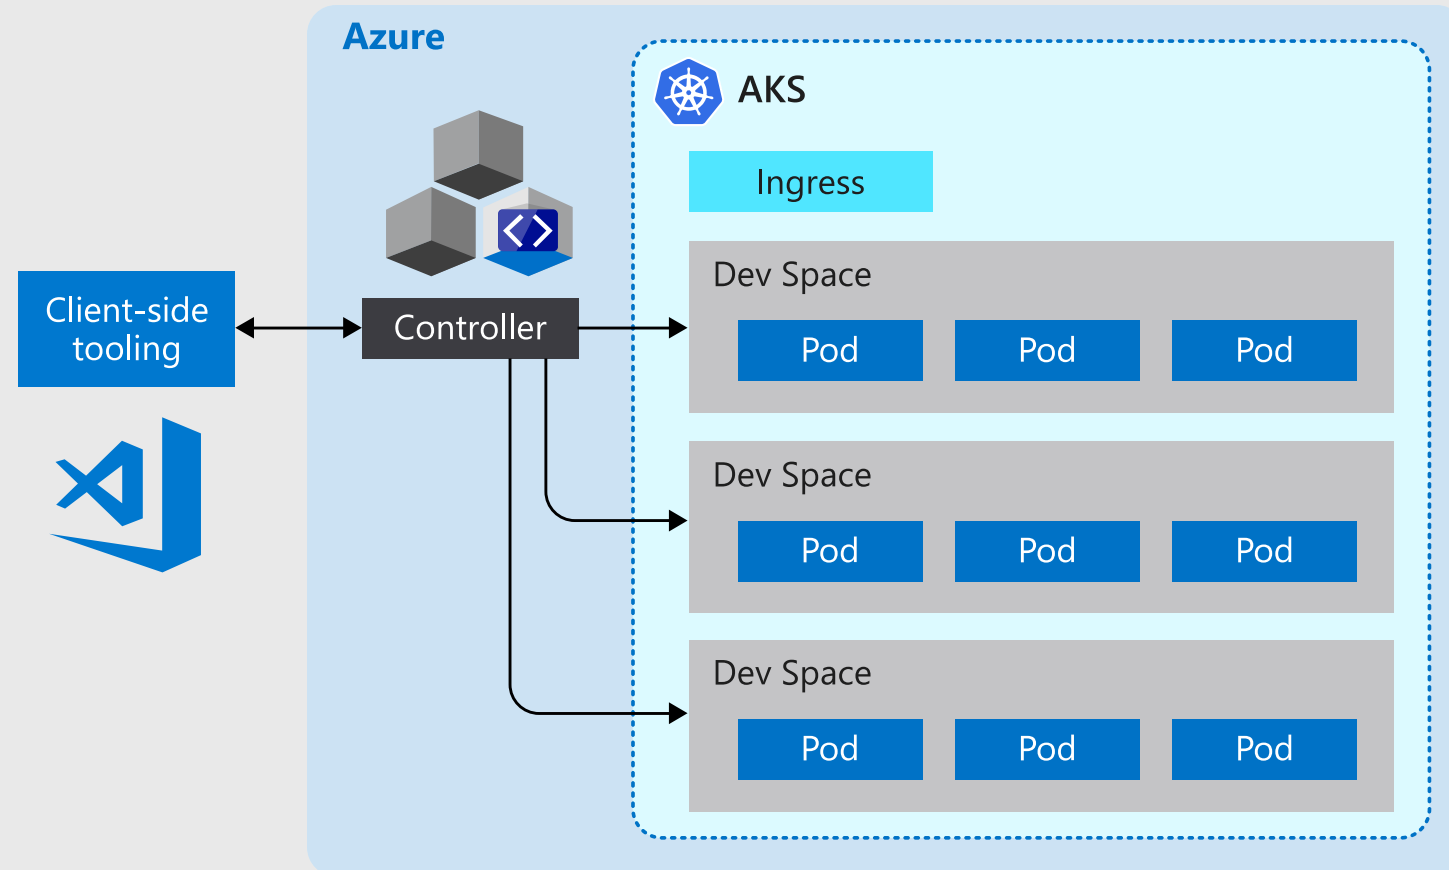
PREVIEW

Minimize local dev machine setup  
and work directly in AKS

Rapidly iterate & debug directly in  
AKS using Visual Studio 2017 or Visual  
Studio Code.

Generate Docker and Kubernetes  
configuration-as-code assets

Share a Kubernetes cluster with your  
team



## Additional Resources

Kubernetes  
Docs

API Reference

Kubernetes  
Hands On Lab

Kubernetes  
Workshop

Katacoda  
Courses

Udemy  
Course



THANK YOU



# Change Log

- 1.0 – Initial release
- 1.1 – Typos and fixes, added final additional resources slide
- 1.2 – Added AKS section & taints/affinity
- 1.3 – Extra sections on AKS
- 1.4 – Minor refresh, ConfigMaps and service mesh